Sharif University of Technology

Department of Mechanical Engineering

Physics-Informed Modeling of Fluid–Particle Transport in Channels

From Dilute Particle Tracking to Gaussian Parcels and Constrictions

Mohammad Reza Sayah

Master's Thesis

Master of Science

Supervisor: Mohammad Ali Bijarchi

Tehran, Iran

October 2026

Abstract

Predicting particle transport requires an accurate carrier flow, a consistent interphase force model, and an evaluation that distinguishes field agreement from reliable particle motion. I develop a staged physics-informed framework that begins with dilute particle tracking in a curved channel, introduces Gaussian parcel coupling in a straight channel, and formulates an extension to a smoothly constricted channel. The baseline combines a continuous velocity–pressure network with a neural approximation of the prescribed Stokes-drag law. Independent finite-element and analytical comparisons give sample-relative errors of 1.510% in curved-channel velocity and 0.863% in gauge-aligned pressure. The normalized drag error is 0.717%. The Gaussian formulation uses the same normalized kernel for fluid interpolation and reaction-force spreading, yielding exact discrete force and work identities on the transfer quadrature. Reduced-model tests recover Stokes relaxation, two-way momentum exchange, Gaussian overlap, and Taylor–Green decay. For these tests, the reported relative errors are below 0.25%; this bound concerns the listed solution observables, not every residual or every channel prediction. The fully coupled Gaussian channel calculations do not yet satisfy all residual criteria and are excluded from the set of accepted predictive results. Likewise, the baseline particle endpoint comparisons identify a remaining near-wall accuracy limitation. The final stage defines a smooth constriction, consistent wall treatment, and a reference-based evaluation protocol; no completed constricted-channel result is claimed. The central contribution is a verified foundation for progressing from one-way particle tracking to conservative regularized coupling, together with explicit conditions for extending that foundation to confined geometries.

Keywords: physics-informed neural networks; particle transport; Stokes drag; Gaussian parcels; two-way coupling; channel constriction; numerical verification.

Contents

Abstract i

Contents ii

List of Figures v

List of Tables vi

Abbreviations and Nomenclature vii

1 Introduction 1

1.1 Motivation . . . 1
1.2 Problem statement . . . 1
1.3 Aim and objectives . . . 2
1.4 Scope and physical assumptions . . . 2
1.5 Study sequence and contributions . . . 3
1.6 Thesis organization . . . 3

2 Literature Review 5

2.1 Modeling levels for particle-laden flow . . . 5
2.2 One-way and two-way coupling . . . 5
2.3 Gaussian parcels and represented populations . . . 6
2.4 Physics-informed representations . . . 6
2.5 Independent references and numerical credibility . . . 7
2.6 Transport accuracy and confinement . . . 8
2.7 Position of the present study . . . 8

3 Mathematical and Evaluation Framework 9

3.1 Baseline nondimensionalization . . . 9
3.2 Geometry and boundary-value matching . . . 9
3.3 Gaussian-channel scales . . . 10
3.4 What is learned . . . 11
3.5 Error measures and pressure gauge . . . 11

3.6 Trajectory evaluation . . . 12
3.7 Evidence hierarchy . . . 12

4 Baseline Particle Problem 13
4.1 Physical problem and parameter set . . . 13
4.2 Flow and force representations . . . 13
4.3 Independent finite-element reference . . . 14
4.4 Analytical plane-flow verification . . . 16
4.5 Curved-channel carrier-flow results . . . 16
4.6 Held-out drag and field comparisons . . . 19
4.7 Seed variation, reference sensitivity, and flux . . . 19
4.8 Particle-transport interpretation . . . 21
4.9 Chapter conclusions . . . 21

5 Gaussian Parcels in a Straight Channel 22
5.1 From individual particles to represented populations . . . 22
5.2 Coupled equations . . . 22
5.3 Normalized Gaussian transfer . . . 23
5.4 Neural representation and hard constraints . . . 24
5.5 Reduced Stokes-response verification . . . 24
5.6 Two-way momentum-exchange verification . . . 26
5.7 Kernel and carrier-flow verification . . . 27
5.8 Status of the coupled straight-channel calculations . . . 28
5.9 Chapter conclusions . . . 30

6 Extension to a Constricted Channel 31
6.1 Purpose and status of the extension . . . 31
6.2 Parameterized geometry . . . 31
6.3 Carrier-flow boundary conditions . . . 32
6.4 Geometry-aware representation . . . 33
6.5 Gaussian support near the narrowing . . . 33
6.6 Particle and parcel wall conditions . . . 34
6.7 Controlled evaluation sequence . . . 35
6.8 Observables and acceptance criteria . . . 36
6.9 Expected mechanisms and testable statements . . . 37
6.10 Chapter conclusions . . . 37

7 Discussion **38**
7.1 What the baseline establishes . . . 38
7.2 What conservative transfer establishes . . . 38
7.3 Optimization, residuals, and independent error . . . 39
7.4 Regularization and model validity . . . 39
7.5 Transferability across the three stages . . . 40
7.6 Implications for the next computational study . . . 40

8 Conclusions and Future Work **42**
8.1 Principal conclusions . . . 42
8.2 Contributions . . . 43
8.3 Priorities for future work . . . 43

References **44**

A Derivations and Transport Limits **46**
A.1 Uniform two-way relaxation . . . 46
A.2 Adjoint transfer and kinetic energy . . . 46
A.3 A limiting pressure-drop estimate . . . 47
A.4 Baseline trajectory criteria . . . 47
A.5 Planar mass convention . . . 48

B Computational and Reproducibility Record **49**
B.1 Completed numerical records . . . 49
B.2 Training and evaluation distinction . . . 49
B.3 Software and execution scope . . . 50
B.4 Figures and numerical transformations . . . 50
B.5 Data and availability record . . . 51

List of Figures

1.1 Study progression and proposed constriction 4

4.1 Baseline flow and force network architecture 15
4.2 Analytical plane Poiseuille verification 17
4.3 Curved-channel fields and discrepancies 18
4.4 Independent field and normalized drag comparisons 19
4.5 Reference-mesh and training-seed sensitivity 20
4.6 Section flux in the curved baseline 20

5.1 Coupled fluid and Gaussian parcel representation 25
5.2 V01: reduced Stokes response 26
5.3 V02: uniform two-way relaxation 27
5.4 Gaussian overlap and bounded-domain normalization 28
5.5 V10: Taylor–Green carrier-flow verification 29

6.1 Proposed constriction geometry and height profiles 32
6.2 Illustration of kernel support near the throat wall 34

List of Tables

2.1 Model classes and the evidence they require. 8

4.1 Reference-mesh assessment. Velocity differences are measured relative to the finest field on the comparison samples. 16
4.2 Supported baseline numerical comparisons. The relative errors refer to their stated sample sets and normalizations. 17
4.3 Field errors for the two completed training seeds. 20

5.1 Accepted solution-observable errors in the reduced Gaussian-model verification problems. Residual criteria are checked separately. 29
5.2 Coupled straight-channel screening. Every required held-out residual must be below 0.03; the momentum residual exceeds that bound in all four cases. These are diagnostic calculations, not accepted transport predictions. 29

6.1 Proposed acceptance protocol for the constricted study. Limits must be fixed before interpreting the associated parameter sweep; these checks have not been executed. 36

7.1 Supported conclusions and their boundaries. 40

A.1 Complete baseline particle comparison; no initial position is removed to improve the aggregate outcome. 48

Abbreviations and Nomenclature

Abbreviation	Meaning
AD	Automatic differentiation
CFD	Computational fluid dynamics
ERPP	Exact regularized point-particle method
FEM	Finite-element method
MLP	Multilayer perceptron
ODE / PDE	Ordinary / partial differential equation
PINN	Physics-informed neural network
RMS	Root-mean-square

Symbol	Meaning
$\boldsymbol{u}, p$	Dimensionless fluid velocity and pressure
$\boldsymbol{q}_i, \boldsymbol{v}_i$	Parcel center and parcel velocity
$\boldsymbol{U}_i$	Fluid velocity interpolated to parcel i
$\boldsymbol{F}_i, \boldsymbol{S}_a$	Represented hydrodynamic drag and reaction source
L_0, U_0, H, B	Reference length, speed, channel height and spanwise depth
$b = B/H$	Reference-depth ratio for planar sphere parcels
d_p, δ_p	Particle diameter scaled by L_0 or H, respectively
$R = \rho_p/\rho_f$	Particle-to-fluid density ratio
$\mathrm{Re}, \mathrm{Re}_p, \mathrm{St}$	Flow Reynolds, slip Reynolds and Stokes numbers
β, α	Mass loading and effective local solid volume fraction
N_i, m_i	Parcel multiplicity and represented dimensionless mass
σ, W_{ai}, Z_i	Gaussian width, normalized transfer weight and normalizer
ΔA_a	Transfer-quadrature area associated with point a
s, η	Centerline coordinate and normalized transverse coordinate
Q, h, H_t	Volume flux per reference depth, half-height and throat height
δ, ℓ, x_c	Constriction severity, half-length and center
E_u, E_p	Discrete relative velocity and gauge-aligned pressure errors

Pressure uses a viscous scale in the baseline and an inertial scale in the Gaussian channel. Each formulation states its scale explicitly. The Gaussian width σ is distinct from both the particle radius and the channel half-height h.

Chapter 1

Introduction

1.1 Motivation

Particle-laden channel flows connect local hydrodynamics with transport outcomes. A particle can follow the carrier flow closely, cross a streamline because of inertia, or approach a wall closely enough to change its final destination. In a narrow passage, these outcomes depend on both the flow field and the force closure. A small average velocity error can coexist with a significant error in a particle's residence time or contact location. Predictive modeling therefore requires more than visually plausible streamlines.

Conventional computational methods provide established ways to discretize the incompressible Navier–Stokes equations and integrate particle motion [6, 12]. Physics-informed neural networks offer a different representation: a differentiable network approximates the solution, while differential equations and boundary conditions constrain its parameters [1, 2]. This representation is attractive when smooth field evaluation and derivatives are needed at many moving particle positions. Its utility must nevertheless be established by comparison with an independent reference and by checks that are sensitive to the intended transport quantities.

I organize this study around a progressive increase in physical coupling and geometric confinement. The first stage considers dilute particles whose motion does not alter the carrier flow. The second introduces represented particle populations and a spatially regularized reaction force in a straight channel. The third defines a smooth constriction as the next geometry to be studied. This progression isolates the consequences of adding two-way coupling before introducing a varying channel height.

1.2 Problem statement

The central problem is to determine which parts of a physics-informed fluid–particle model have sufficient numerical support for transport prediction. There are three distinct questions. The first concerns the carrier field: does the predicted velocity and pressure agree with a solution of the same boundary-value problem? The second concerns the force transfer: does the discrete coupling apply equal and opposite hydrodynamic forces without introducing artificial interphase power? The third

concerns particle transport: are trajectories, contact events, and transit times accurate for the geometry under consideration?

These questions are related but are not interchangeable. An algebraically conservative transfer operator does not eliminate the residual error of a trained fluid field. A low drag-network error does not correct an inaccurate near-wall velocity. A successful straight-channel benchmark does not establish performance at a constriction. I therefore treat each claim at the level at which it has been evaluated.

1.3 Aim and objectives

The aim is to establish a physically consistent and numerically assessable framework for physics-informed particle transport in channels. The study has four objectives:

1. Formulate and evaluate a one-way coupled baseline with independent carrier-flow and analytical drag references.
2. Derive a Gaussian parcel coupling that preserves force and interphase work on its transfer quadrature.
3. Evaluate the components of the coupled formulation using independent analytical and numerical verification problems.
4. Specify a constricted-channel extension whose geometry, boundary conditions, and acceptance criteria can be tested without altering the meaning of the earlier benchmarks.

The corresponding research questions are: How accurately does the baseline recover independently computed channel flow? Which conservation properties follow from the Gaussian discretization, and which require additional solution accuracy? What evidence is still needed before the formulation can support predictions in a constricted channel?

1.4 Scope and physical assumptions

The completed channel models describe incompressible Newtonian flow and spherical heavy particles moving in a plane. The baseline uses one-way Stokes drag and a prescribed irreversible wall-sticking rule. The Gaussian model represents spherical-particle populations through regularized force distributions. The Gaussian cloud is a force-transfer representation; it is not a resolved solid sphere and it is not an independently validated description of a dense particle agglomerate.

The two-dimensional carrier does not resolve spanwise motion. No inference about a finite-depth three-dimensional channel follows from an extruded view. Particle

rotation, lubrication, history force, and hydrodynamic interactions of resolved finite-size bodies lie outside the accepted baseline. The constriction is considered as a static wall geometry rather than a dynamically growing deposit or a biological vessel.

1.5 Study sequence and contributions

Figure 1.1 summarizes the study sequence. The first contribution is a baseline evaluation that combines analytical plane flow, an independently assembled finite-element reference, a held-out drag test, and an explicit distinction between field and trajectory accuracy. The second is a consistent planar sphere-parcel formulation with shared Gaussian interpolation and spreading. The third is a verification chain connecting reduced relaxation problems, normalization identities, and an unsteady incompressible-flow benchmark. The final contribution is a defined constricted-channel study design that preserves the same requirements for reference matching and conservation.

Completed findings are presented for the baseline carrier field and the Gaussian model's verification problems. The coupled Gaussian channel predictions have unresolved residual errors. The constricted-channel formulation is prospective. This division determines the interpretation of the conclusions and prevents a proposed geometry from being presented as an executed calculation.

1.6 Thesis organization

Chapter 2 reviews the modeling choices and the requirements for credible numerical evidence. Chapter 3 defines the common nondimensional and evaluation framework. Chapter 4 develops the baseline particle problem and reports its supported results. Chapter 5 introduces the straight-channel Gaussian formulation and its verification results. Chapter 6 defines the constricted channel and the proposed evaluation protocol. Chapter 7 compares the implications of the stages, and chapter 8 gives the conclusions and priorities for further work. Supplementary derivations and the computational record appear in the appendices.

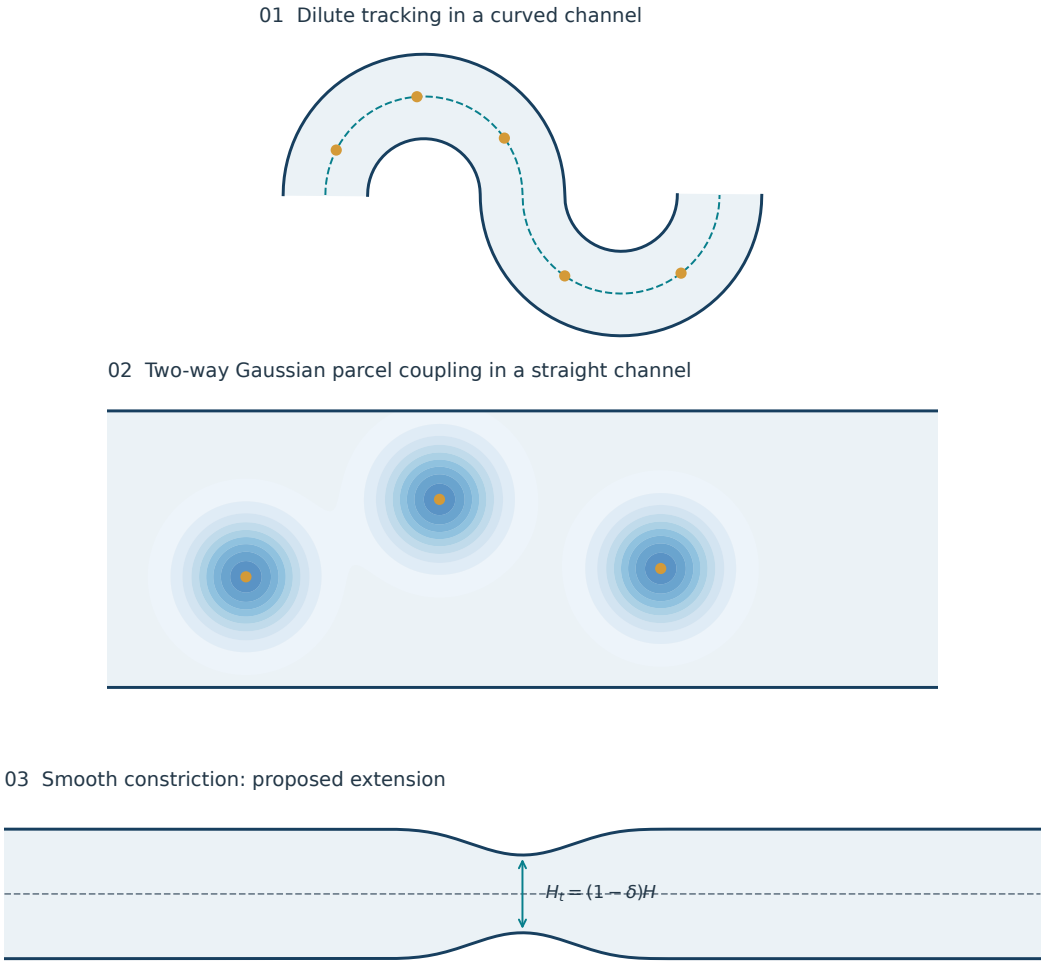


Figure 1.1. Study progression. The first two panels identify the model classes investigated; the third defines the proposed constricted-channel extension. All three panels are geometric illustrations, not simulation fields. Gaussian shading depicts kernel support rather than a measured particle concentration.

Chapter 2

Literature Review

2.1 Modeling levels for particle-laden flow

The hydrodynamics of a particle suspension can be described at different levels of resolution. A particle-resolved method imposes boundary conditions on each particle surface. It can represent the disturbance flow, finite-size effects, and surface stresses when these are sufficiently resolved. Direct forcing immersed-boundary methods provide one established route to such calculations [8]. Their additional resolution requirements are justified when particle-scale flow is itself an essential part of the problem.

A point-particle model instead advances particle centers using a force law obtained under specified physical assumptions. The classical equation of motion for a small rigid sphere includes effects beyond steady drag [5]. The heavy-particle, low-slip-Reynolds-number approximation used here retains a much narrower description. It is appropriate to assess this approximation as a simplified model rather than attribute to it the full physics of the parent equation.

The strength of a point-particle approach is its economy of representation. Its limits follow from the same reduction. A particle center and diameter do not impose no-slip conditions on an actual moving surface. Finite-size wall stresses, rotational dynamics, and lubrication do not appear unless additional closures are included. These distinctions are particularly important when a channel becomes narrow relative to particle diameter [6].

2.2 One-way and two-way coupling

In one-way coupling, the fluid field determines the particle force but the particle reaction is neglected in the fluid equations. This separation is useful for studying dilute transport because the fluid solution can be evaluated independently of particle initial conditions. Small particle volume fraction is a necessary consideration, although it does not alone guarantee negligible feedback when the particle density is high. Mass loading and the relevant time scales also matter.

Two-way coupling includes the hydrodynamic reaction in the carrier-fluid momentum balance. The coupling must specify how a particle samples the fluid and how

its reaction is distributed to the Eulerian field. Treating either operation independently can generate an inconsistent power exchange even when the integrated force is correct. A shared transfer operator provides a direct means to avoid this particular inconsistency.

Regularized point-particle methods provide a principled context for this transfer. In the exact regularized point-particle method, the treatment of the disturbance and its regularization is tied to the governing flow equations [7]. A generic normalized Gaussian deposition does not inherit all of that method's properties. In particular, sampling the disturbed velocity at a force source is not equivalent to recovering the undisturbed velocity required by an isolated-particle drag law.

2.3 Gaussian parcels and represented populations

A parcel represents a population of physical particles through a multiplicity. Its represented mass must scale with that multiplicity, and its reaction force must use the same convention. A Gaussian kernel spreads that force over a finite region, avoiding a singular source in the carrier equations. The width of the kernel is a modeling parameter with a direct effect on the spatial distribution of forcing.

Three lengths must be kept distinct: the physical diameter, the numerical or model regularization width, and the local channel height. Equating the Gaussian width with a particle radius without justification would blur point-particle and finite-size interpretations. Likewise, changing the width while keeping quadrature resolution fixed does not constitute an ordinary grid-convergence study: it changes the regularized model being solved.

Near a wall, a free-space Gaussian is truncated by the fluid domain. A discrete normalization can preserve its integrated weight on the chosen quadrature. That normalization establishes a transfer identity; it does not derive a wall-corrected drag law. This distinction carries directly into the constricted-channel extension developed later.

2.4 Physics-informed representations

Physics-informed neural networks approximate solution fields with differentiable functions and train those functions using equation and boundary residuals [1]. They provide continuous field values without a separate spatial interpolant. Automatic differentiation also allows the same representation to supply the derivatives entering the equations. These features are useful for particle transport, where the evaluation locations move continuously.

Physics information in the objective is not a guarantee that optimization finds a sufficiently accurate solution. Residual terms can have different scales and different optimization behavior. Gradient imbalance and training difficulties are documented limitations of PINNs [3, 4]. A weighted objective may decrease while a transport-sensitive region remains inaccurate. Consequently, reporting only the optimization history would provide insufficient support for the predictions considered here.

Hard constraints offer an additional design choice. An ansatz can impose an initial condition, periodicity, or homogeneous wall velocity exactly, leaving the network to represent a correction that respects those conditions. This reduces the number of requirements enforced through a penalty. It does not remove the need to evaluate momentum, continuity, and particle-motion residuals at independent points.

The baseline and Gaussian stages use different neural roles. The baseline flow network approximates a steady velocity–pressure field, whereas the force network approximates a prescribed algebraic law. In the Gaussian stage, separate fluid and parcel networks are coupled through common residuals. This is a continuous-time coupled forward representation rather than a trajectory regression against reference data.

2.5 Independent references and numerical credibility

The finite-element method supplies an independently discretized reference for the baseline carrier flow. Taylor–Hood velocity–pressure elements are a standard mixed formulation for incompressible flow [9]. Their assembly here uses `scikit-fem` [10]. An independent discretization is valuable because its numerical errors differ from those of the neural representation. It must still be evaluated for nonlinear convergence, incompressibility, and mesh sensitivity.

Verification and physical validation address different questions [11]. Agreement with an analytical solution tests whether a specified mathematical problem is solved correctly. Agreement with another numerical solver tests consistency for a matched model and boundary-value problem. Experimental validation additionally requires suitable measurements, controlled conditions, and compatible physical assumptions. The results in this thesis primarily concern the first two questions.

A relative velocity error should be interpreted together with local discrepancies, boundary errors, flux variation, and particle-motion diagnostics. Pressure requires particular care because incompressible velocity boundary conditions determine pressure only up to an additive constant. Comparing unaligned pressure arrays can produce a large error without a corresponding physical disagreement. Conversely, an error measured on a section where the reference pressure is constant can have an ill-

conditioned denominator. A spatially extended, gauge-aligned comparison avoids these problems.

2.6 Transport accuracy and confinement

Particle endpoints depend on a time-integrated response to the carrier field. A small discrepancy may accumulate along a long path or determine whether a particle approaches a capture boundary. Near-wall outcomes can therefore be less accurate than bulk flow metrics suggest. A meaningful trajectory comparison must pair the same initial conditions, restrict its position error to the common time interval, and separately report whether the two solutions reach the same endpoint class.

Confinement adds a physical limitation to this numerical sensitivity. When the local channel height approaches the particle diameter, the unbounded Stokes law becomes a questionable description even if the numerical trajectory is well resolved. A study of narrowing should therefore track both numerical agreement and the changing ratio of particle diameter to local gap. A dilute regularized model cannot be used to infer dense blockage merely because a narrow geometry is drawn.

2.7 Position of the present study

Table 2.1 summarizes the modeling levels relevant to the study. The purpose is to build a verified chain across the first two levels and specify the requirements for a constrained extension. The focus is neither a new empirical drag correlation nor a replacement for resolved particle hydrodynamics. The main methodological question is whether a neural representation can support the stated fluid and particle quantities while preserving a consistent transfer of force and work.

Table 2.1. Model classes and the evidence they require.

Model	Main representation	Essential assessment
One-way point particles	Carrier field and particle ODEs	Field error, slip regime, trajectory agreement
Gaussian parcels	Population-weighted regularized forcing	Shared transfer, residuals, budgets, width sensitivity
Resolved particles	Moving solid surfaces and surrounding flow	Surface resolution, forces, torques and gap flow
Constricted dilute model	Variable-height fluid domain with particles	Wall geometry, flux, throat resolution and closure validity

Chapter 3

Mathematical and Evaluation Framework

3.1 Baseline nondimensionalization

Let L_0 and U_0 denote reference length and velocity, and let ρ_f and μ denote the fluid density and dynamic viscosity. The baseline uses

$$\boldsymbol{x}^* = \frac{\boldsymbol{x}}{L_0}, \quad t^* = \frac{tU_0}{L_0}, \quad \boldsymbol{u}^* = \frac{\boldsymbol{u}}{U_0}, \quad p_v^* = \frac{pL_0}{\mu U_0}, \quad \mathrm{Re} = \frac{\rho_f U_0 L_0}{\mu}. \tag{3.1}$$

Stars are suppressed after the scales have been specified. The steady carrier equations are

$$\nabla \cdot \boldsymbol{u} = 0, \qquad \mathrm{Re}(\boldsymbol{u} \cdot \nabla)\boldsymbol{u} = -\nabla p_v + \nabla^2 \boldsymbol{u}. \tag{3.2}$$

Using a viscous pressure scale places Re on the convective term. Pressure values from this model cannot be compared directly with a model using an inertial pressure scale without conversion.

For a spherical heavy particle, the retained force is the creeping-flow drag. With $R = \rho_p/\rho_f$ and $d_p = d/L_0$, its dimensionless equations are

$$\frac{d\boldsymbol{q}}{dt} = \boldsymbol{v}, \qquad \frac{d\boldsymbol{v}}{dt} = \frac{18}{R\mathrm{Re}d_p^2}[\boldsymbol{u}(\boldsymbol{q}) - \boldsymbol{v}] = \frac{\boldsymbol{u}(\boldsymbol{q}) - \boldsymbol{v}}{\mathrm{St}}, \qquad \mathrm{St} = \frac{R\mathrm{Re}d_p^2}{18}. \tag{3.3}$$

The slip Reynolds number is

$$\mathrm{Re}_p = \mathrm{Re}d_p|\boldsymbol{u}(\boldsymbol{q}) - \boldsymbol{v}|. \tag{3.4}$$

The physical rationale for Stokes drag requires $\mathrm{Re}_p \ll 1$, along with a small particle relative to the spatial scales of the carrier flow.

3.2 Geometry and boundary-value matching

The baseline channel is a constant-width tube around a planar centerline $\boldsymbol{c}(s)$. Let $\boldsymbol{n}(s)$ denote its unit normal and let h denote the half-width. Coordinates are mapped

by

$$\boldsymbol{x}(s,\eta) = \boldsymbol{c}(s) + h\eta\boldsymbol{n}(s), \qquad 0 \le s \le L, \quad -1 \le \eta \le 1. \tag{3.5}$$

The carrier velocity vanishes on both walls. At each end section, a parabolic velocity in the local tangent direction is prescribed:

$$\boldsymbol{u}(s_{\mathrm{end}},\eta) = (1-\eta^2)\boldsymbol{t}(s_{\mathrm{end}}). \tag{3.6}$$

Thus the reference problem uses velocity conditions at both ends, rather than an imposed pressure outlet. A pressure gauge fixes the remaining constant. The neural and finite-element comparisons must use this same problem.

The associated volume flux per unit reference depth is

$$Q_0 = h\int_{-1}^{1}(1-\eta^2)\,d\eta = \frac{4h}{3}. \tag{3.7}$$

For a steady incompressible flow without sources of fluid mass, the section flux should be independent of s. This makes flux variation a useful integrated diagnostic alongside local differential residuals.

3.3 Gaussian-channel scales

For the Gaussian stage, the reference length is the channel height H. The pressure scale is inertial:

$$p_i^* = \frac{p}{\rho_f U_0^2}, \qquad \mathrm{Re}_H = \frac{U_0 H}{\nu}, \qquad \delta_p = \frac{d}{H}. \tag{3.8}$$

For identical length and velocity scales, the two pressure conventions satisfy $p_v^* = \mathrm{Re}_H p_i^*$. The governing equation is correspondingly written with $1/\mathrm{Re}_H$ multiplying viscous diffusion.

The carrier is two-dimensional, but the represented physical particles are spheres. A reference depth B is therefore required to convert represented sphere mass into a planar source. With $b = B/H$ and multiplicity N_i, the dimensionless represented mass is

$$m_i = \frac{N_i R\pi\delta_{p,i}^3}{6b}, \qquad \mathrm{St}_i = \frac{R\mathrm{Re}_H\delta_{p,i}^2}{18}, \qquad \beta = \frac{\sum_i m_i}{|\Omega|}. \tag{3.9}$$

This convention scales physical mass by $\rho_f H^2 B$. The kernel has units of inverse area before nondimensionalization. Replacing it with a three-dimensional volume kernel without changing the mass convention would be dimensionally inconsistent.

3.4 What is learned

The neural representations approximate fields or trajectories; the physical equations specify their constraints. The baseline force network learns a known drag response from an algebraic residual. It does not infer an unknown force law from experiments. The Gaussian parcel network predicts time-dependent position and velocity corrections whose consistency is enforced by both kinematic and dynamic residuals.

Adam is used for initial optimization, with optional limited-memory quasi-Newton refinement [14, 15]. The optimizer sequence does not define the credibility of the final state. Checkpoint selection uses a fixed residual monitor, whereas the reported field comparison uses an independent reference. A longer training budget is a computational choice, not an accuracy certificate.

3.5 Error measures and pressure gauge

For values at M independent locations, the sample-relative velocity error is

$$E_u = \left[\frac{\sum_{j=1}^{M} |\boldsymbol{u}_j^{\mathrm{PINN}} - \boldsymbol{u}_j^{\mathrm{ref}}|^2}{\sum_{j=1}^{M} |\boldsymbol{u}_j^{\mathrm{ref}}|^2} \right]^{1/2}. \tag{3.10}$$

These are discrete norms over the specified samples. Unless area weights are explicitly included, they are not exact domain-integral norms. For pressure, I first remove the arithmetic mean of each compared array:

$$\widetilde{p}_j = p_j - \frac{1}{M} \sum_{k=1}^{M} p_k, \qquad E_p = \frac{\|\widetilde{\boldsymbol{p}}^{\mathrm{PINN}} - \widetilde{\boldsymbol{p}}^{\mathrm{ref}}\|_2}{\|\widetilde{\boldsymbol{p}}^{\mathrm{ref}}\|_2}. \tag{3.11}$$

The pressure samples span the streamwise direction so that the reference denominator is meaningful. Mean subtraction removes only a constant; it does not rescale a pressure gradient or modify the predicted dynamics.

The section-flux diagnostic is

$$E_Q = \max_s \left| \frac{Q(s)}{Q_0} - 1 \right|. \tag{3.12}$$

Reported differential residuals use their stated nondimensional equations. Their magnitudes should not be compared across pressure scales, equation normalizations, or geometries without accounting for those differences.

3.6 Trajectory evaluation

Consider a neural and reference trajectory with the same initial condition. If their common interval is $[0, T_c]$, the position discrepancy is measured over that interval and normalized by channel width $W = 2h$. Endpoint agreement is assessed separately using the terminal event class. I also retain the fraction of the longer trajectory covered by the common interval. A low error calculated over a short common interval is not sufficient evidence of agreement over the full transit.

Integration success has a narrower meaning: it indicates that the numerical ODE solver completed its assigned calculation without reporting failure. It does not establish that the neural velocity or force was accurate. Accordingly, the acceptance criteria distinguish integration status, position discrepancy, common-time coverage, and endpoint agreement.

3.7 Evidence hierarchy

I use four levels of numerical evidence. Algebraic identities establish properties of the coupling operators. Analytical benchmarks test reduced equations with known solutions. Independently discretized references test the field approximation for a matched problem. Finally, transport comparisons assess the accumulated consequences of the field and force errors. None of these alone establishes experimental agreement.

For the baseline, the numerical screening limits are 5% for field and normalized drag errors, 2% for boundary RMS and reference-interpolation error, and 5% for maximum flux deviation. Baseline particle acceptance is stricter in its combination of conditions: every tested endpoint class must agree, every width-normalized RMS position error must be below 5%, and every common-time fraction must exceed 95%. For the Gaussian verification tests, the stated relative-error and independent residual criteria are 0.03. These limits are study criteria rather than universal physical thresholds. Their actual outcomes are reported where the corresponding results are presented.

Chapter 4

Baseline Particle Problem

4.1 Physical problem and parameter set

The baseline separates carrier-flow approximation from particle response. The carrier is steady and two-dimensional. Particles are spherical, heavy, and dilute, and their forces do not feed back into the fluid. The principal channel consists of two connected semicircular bends with nominal inner and outer radii $R_{\text{in}} = 0.2$ and $R_{\text{out}} = 0.5$. Its width is $W = 0.3$ and its centerline length is approximately 2.1991, all scaled by L_0. The Reynolds number is $\text{Re} = 1$.

The baseline particle parameters are $R = 1000$ and $d_p = 0.025$. The resulting Stokes number is $\text{St} \simeq 0.03472$. The channel wall rule stops the particle when its center approaches within one radius of a wall. This is a specified irreversible sticking model, not a prediction of adhesion chemistry or a resolved impact process. In the full trajectory set, the largest recorded slip Reynolds number is approximately 0.01221, which supports the low-slip-Reynolds-number assumption for these calculations. The ratio $d_p/W \simeq 0.0833$ still warrants caution for near-wall motion.

4.2 Flow and force representations

4.2.1 FlowNet

The flow network maps five input slots to velocity and pressure:

$$(x, y, R_{\text{in}}, R_{\text{out}}, \text{Re}) \longmapsto (u, v, p_v). \tag{4.1}$$

It has five hidden layers with 70 tanh units per layer. The spatial inputs are centered and scaled, and the pressure output is scaled with a channel-based pressure magnitude. The geometry and Reynolds inputs are fixed within an individual trained model. Their presence does not establish generalization across untrained geometries or Reynolds numbers.

The objective combines momentum and continuity residuals, wall and end velocity conditions, a pressure gauge, and a section-flux condition. If $r_c = \nabla \cdot \boldsymbol{u}$ and $\boldsymbol{r}_m$ denotes

the residual of equation (3.2), the scaled physics term is

$$\mathcal{L}_{\text{phys}} = \left\langle (hr_c)^2 \right\rangle + \left\langle |h^2 \boldsymbol{r}_m|^2 \right\rangle . \tag{4.2}$$

Here the averages follow the componentwise mean-square convention of the implementation. The complete objective uses weights 20 for the velocity boundary term, 1 for the scaled gauge term, and 5 for the normalized flux term. Finite-element interior values are not used in this objective.

4.2.2 ForceNet

The second network has five hidden layers with 90 tanh units and maps

$$(u, v, p_v, d_p, x_p, y_p, u_p, v_p) \longmapsto (F_x, F_y). \tag{4.3}$$

The force scale is the displaced-fluid mass times U_0^2/L_0. Under this convention, the dimensionless Stokes force and particle acceleration are

$$\boldsymbol{F}^* = \frac{18}{\mathrm{Re} d_p^2}(\boldsymbol{u} - \boldsymbol{v}), \qquad \frac{d\boldsymbol{v}}{dt} = \frac{\boldsymbol{F}^*}{R}. \tag{4.4}$$

The neural approximation is trained on the normalized slip response. Subtracting its evaluation at zero slip from its evaluation at the actual slip enforces exactly zero force when $\boldsymbol{u} = \boldsymbol{v}$. The network therefore approximates a specified drag relation. Its agreement with that relation is an approximation test, not discovery of a new constitutive law.

Figure 4.1 gives a visual summary of both networks. The role of ForceNet is intentionally kept separate from the carrier-field approximation so that the two errors can be evaluated independently.

4.3 Independent finite-element reference

The reference solves the same incompressible equations using quadratic velocity and linear pressure Taylor–Hood elements [9, 10]. A Picard iteration treats the convective term, and sparse direct solution is used for the linear subproblems. The finest channel mesh contains 3,026 triangles, 12,530 velocity degrees of freedom, and 1,620 pressure degrees of freedom.

Three meshes assess sensitivity. Table 4.1 shows that the difference between the two finest velocity fields is 0.173%. The finest reference has a nonlinear residual of approximately 1.60×10^{-10}. Its reported divergence norm is not identically zero;

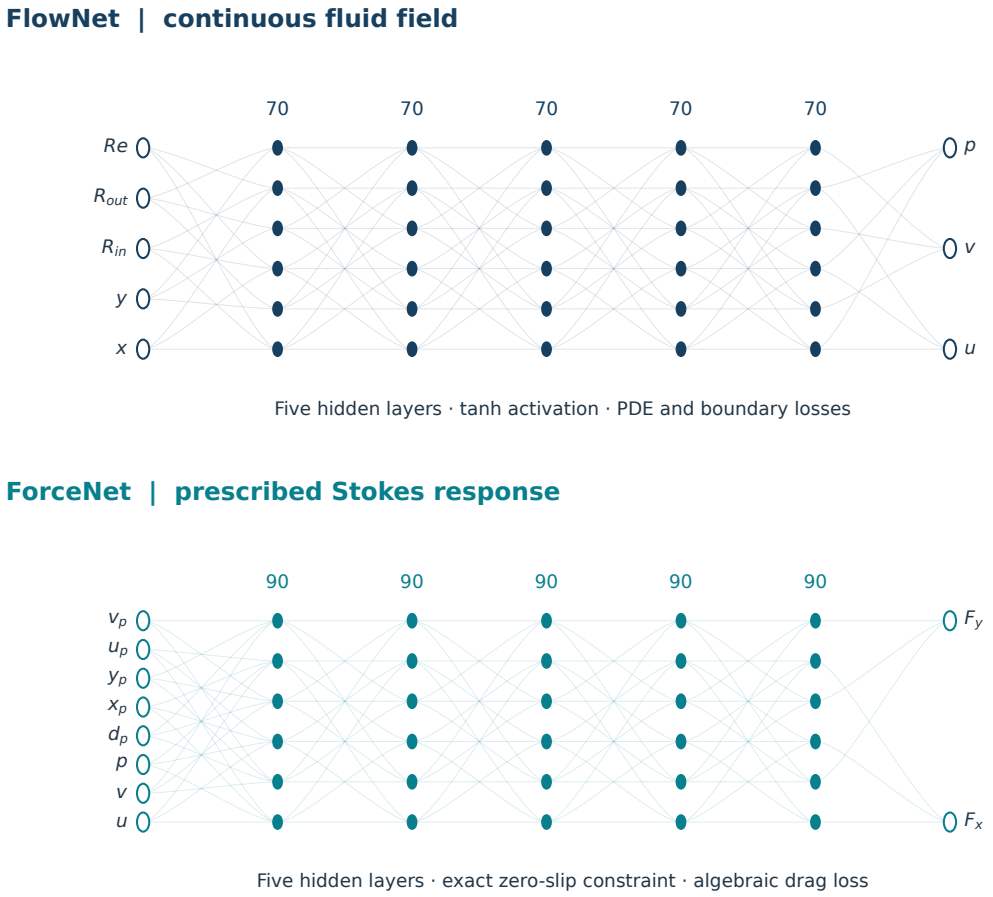


Figure 4.1. Baseline neural architecture. Each row represents a fully connected MLP; the few displayed nodes are illustrative, while the numbers above hidden layers give their actual widths. FlowNet supplies velocity and pressure. ForceNet approximates the normalized Stokes response and uses a zero-slip subtraction. The first network is constrained by PDE and boundary residuals; the second by an algebraic drag residual.

the mixed discretization enforces incompressibility weakly. The mesh differences are sensitivity indicators, not a formal uncertainty estimate or a grid-convergence index [11].

Table 4.1. Reference-mesh assessment. Velocity differences are measured relative to the finest field on the comparison samples.

Triangles	Velocity DOFs	Pressure DOFs	Difference (%)	Nonlinear residual
290	1298	180	1.141	6.735×10^{-11}
1298	5474	720	0.173	1.097×10^{-10}
3026	12530	1620	Reference	1.600×10^{-10}

The first successful training seed is used for the principal comparison; it is not selected by its error against the reference. The reported final evaluation restores completed checkpoints. Consequently, the training budgets associated with those checkpoints, rather than the defaults of a later evaluation command, identify the computations that produced the weights. The computational record in chapter B preserves this distinction.

4.4 Analytical plane-flow verification

Before interpreting curved-channel results, I consider plane Poiseuille flow in a straight channel with half-width h. In the baseline units, the analytical solution is

$$u = 1 - (y/h)^2, \qquad v = 0, \qquad p_v = -\frac{2x}{h^2} + C. \tag{4.5}$$

The finite-element relative velocity error is approximately 8.19×10^{-10}. The PINN velocity error is 0.217%, and its gauge-aligned pressure error is 0.180%. The velocity metric includes both components; figure 4.2 displays the streamwise profile for readability. Pressure is evaluated on points distributed along the channel, with each array's mean removed.

The near-exact finite-element result for this polynomial plane flow is consistent with the approximation spaces and simple geometry. It should not be interpreted as the error of the curved-channel reference. The curved problem requires its own mesh-sensitivity assessment.

4.5 Curved-channel carrier-flow results

Figure 4.3 compares the principal neural field with the finite-element field. Both solutions recover the curved high-speed core and the pressure decrease through the

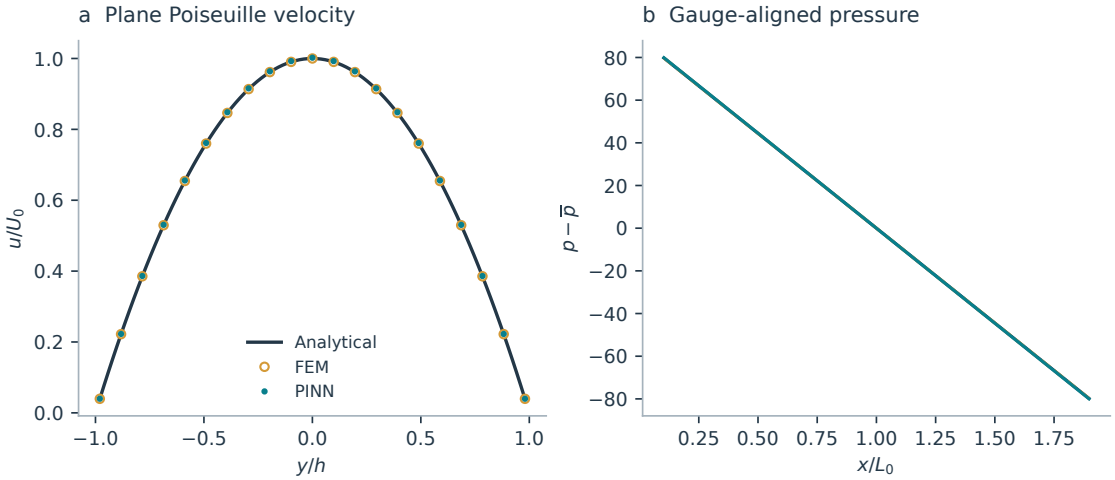


Figure 4.2. Analytical verification of the baseline flow solvers. Panel (a) compares the streamwise velocity profile with the parabolic solution. Panel (b) compares pressure after constant-gauge alignment on spatially distributed samples. Lines and markers retain the stored predictions; no fitted correction is applied.

channel. Common color limits are used for corresponding speed and pressure panels. The discrepancy panels show that agreement is spatially nonuniform; a small global norm should not be interpreted as uniformly small local error.

The velocity error is 1.510%, and the centered pressure error is 0.863% on the independent field-comparison samples. The maximum absolute velocity discrepancy is approximately 0.05401. These values support agreement for the specified boundary-value problem. They do not establish zero equation residual: the separately recorded raw continuity and momentum RMS values are approximately 0.1781 and 1.8296 in the baseline equation convention. Those residuals, the field errors, and the flux diagnostic measure different aspects of the approximation.

Table 4.2. Supported baseline numerical comparisons. The relative errors refer to their stated sample sets and normalizations.

Problem	Quantity	Error
Plane Poiseuille PINN	Velocity	0.217%
Plane Poiseuille PINN	Centered pressure	0.180%
Plane Poiseuille FEM	Velocity	8.192×10^{-10}
Curved-channel PINN	Velocity	1.510%
Curved-channel PINN	Centered pressure	0.863%
ForceNet	Normalized drag	0.717%
Section flux	Maximum relative deviation	2.266%
Two finest FEM meshes	Velocity difference	0.173%

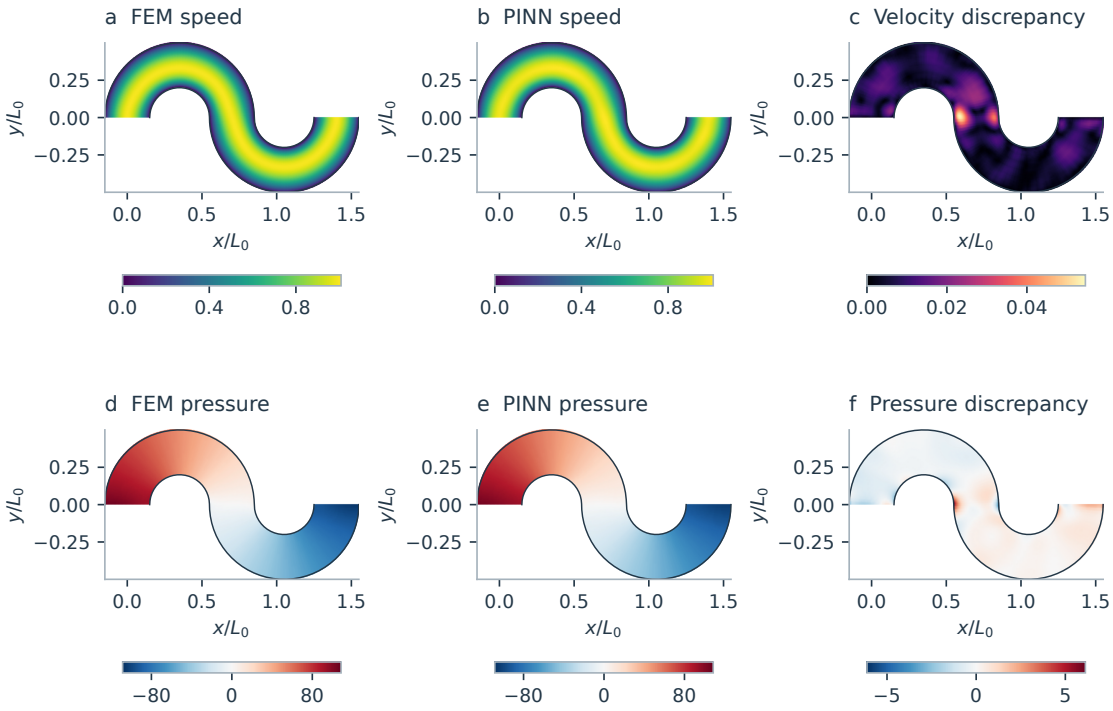


Figure 4.3. Curved-channel carrier field at Re = 1. The speed panels use $|\boldsymbol{u}|/U_0$ and identical scales; the pressure panels use centered viscous-scale pressure and identical scales. Velocity discrepancy is the Euclidean component error, whereas pressure discrepancy is signed. The fields are evaluated on the same stored mapped grid; no spanwise solution is implied.

4.6 Held-out drag and field comparisons

The normalized ForceNet error is 0.717% on the independent 2,048-sample evaluation used for the reported metric. Figure 4.4 presents velocity and pressure parity together with a separate 1,000-sample drag diagnostic. The plotted drag set and the aggregate 2,048-sample test are distinct; they need not have identical numerical errors. This distinction avoids assigning one test's error to another test's points.

The drag comparison supports the network's approximation of the retained creeping-flow law. It cannot establish the accuracy of neglected lift, history, lubrication, or wall corrections. Direct analytical Stokes drag also remains an essential baseline for judging the value of a learned approximation of this known relation.

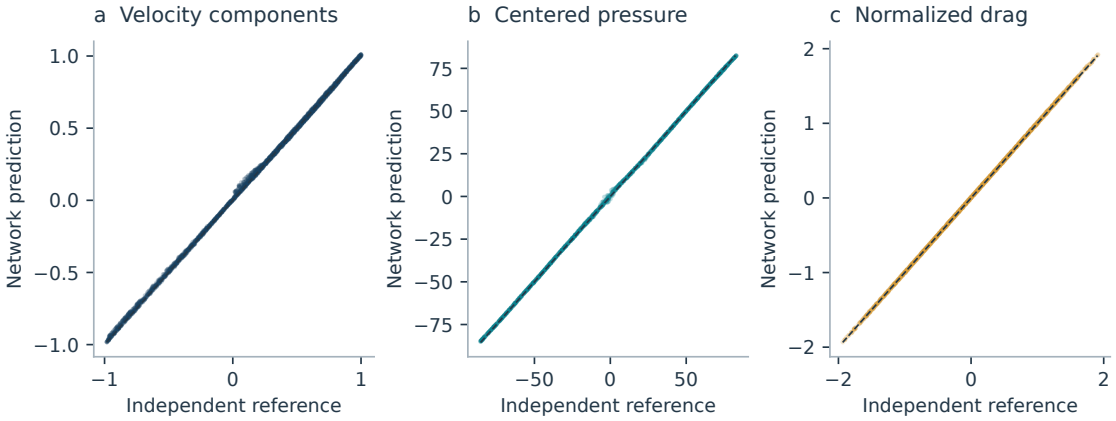


Figure 4.4. Independent numerical comparisons. Velocity components and gauge-aligned pressure are compared with the finite-element reference. The third panel compares normalized learned drag with the specified Stokes-slip response on the stored 1,000-sample diagnostic. The dashed lines indicate equality. The aggregate drag error in table 4.2 comes from a separate 2,048-sample test.

4.7 Seed variation, reference sensitivity, and flux

Two independently trained seeds give velocity errors of 1.510% and 1.429%, and pressure errors of 0.863% and 0.886%, respectively. Their mean velocity error is 1.469%, with a sample standard deviation of 0.057 percentage points. Two seeds are sufficient to reveal some training dependence but are insufficient for a strong statistical uncertainty statement. No confidence interval is inferred.

The maximum section-flux deviation is 2.266%. Figure 4.6 shows the full stored section sequence normalized by the prescribed flux. The variation is finite and should remain visible even though the study's 5% flux criterion is met. Together with boundary RMS of approximately 0.00420, this result provides an integrated check on the velocity

Table 4.3. Field errors for the two completed training seeds.

Seed	Velocity error (%)	Pressure error (%)
42	1.510	0.863
43	1.429	0.886

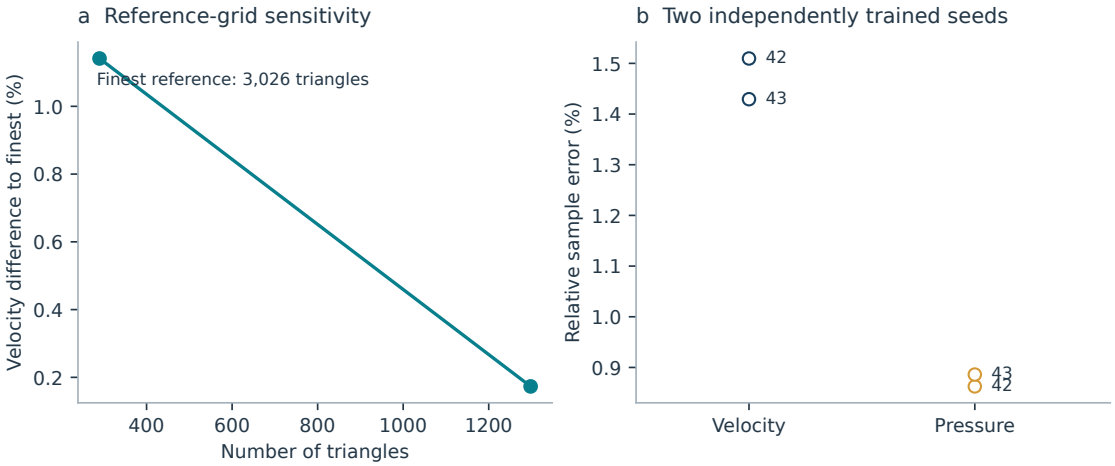


Figure 4.5. Reference and training sensitivity. The left panel compares the two coarser finite-element meshes with the finest reference; the right panel shows the two completed training seeds. The finest mesh is a comparison reference, not an exact solution. Seed labels identify the individual observations rather than confidence bounds.

field that complements the sample-relative error.

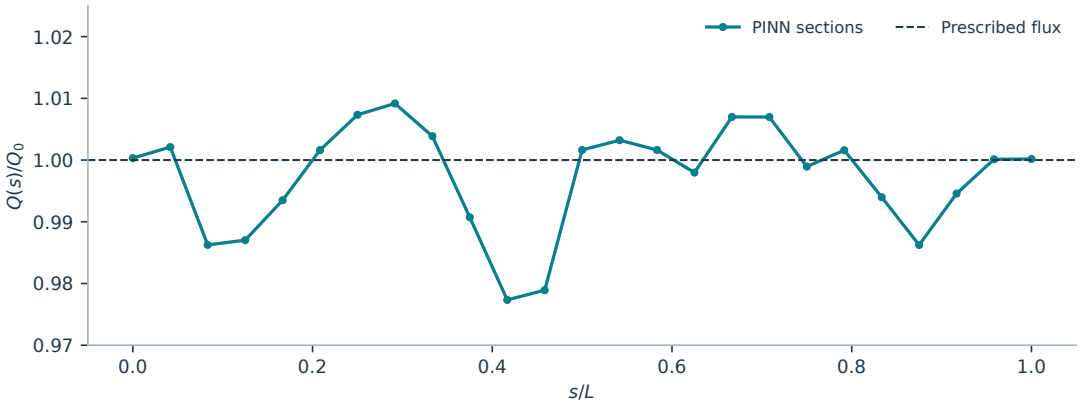


Figure 4.6. Section flux in the curved baseline, normalized by the prescribed value $Q_0 = 4h/3 = 0.2$. The dashed line is the exact constant-flux target, and every stored PINN section is shown. The maximum deviation is 2.266%; the curve has not been rescaled to enforce agreement.

4.8 Particle-transport interpretation

The baseline evaluation completes 119 particle or tracer integrations without a recorded integration failure. The reference velocity interpolator differs from direct finite-element evaluation by 0.177% on its assessment samples. These are useful computational checks, but neither alone proves accurate particle outcomes.

Among the seven baseline particle initial positions, six endpoint classes agree with the reference. The largest width-normalized RMS position discrepancy is approximately 34.97%, and the lowest common-time coverage is approximately 73.55%. The endpoint, position, and time-coverage criteria therefore do not all pass. Detailed values are retained in table A.1 so that the scope of the carrier-flow findings is unambiguous. Particle deposition maps and parameter-sweep trajectories are not used as accepted results in this chapter.

This outcome is physically and numerically informative. Bulk field agreement and a good approximation of a specified drag law are necessary components, but their combination is not enough to certify wall-sensitive transport. The justified baseline conclusion is that the carrier and drag components agree well with their stated references, while the full trajectory prediction remains an improvement target.

4.9 Chapter conclusions

The baseline establishes a reproducible carrier-flow and drag foundation for the subsequent coupling study. Analytical plane flow, independent curved-channel fields, two completed training seeds, and section flux provide complementary evidence. The results also set a boundary on the claim: the present field accuracy does not establish reliable deposition or all baseline particle endpoints. That boundary becomes more important when reaction forcing and geometric narrowing are introduced.

Chapter 5

Gaussian Parcels in a Straight Channel

5.1 From individual particles to represented populations

The second stage adds hydrodynamic feedback while retaining a simple carrier geometry. A parcel represents N_i identical spherical particles whose centers are approximated by a common planar state. Its mass and hydrodynamic reaction are weighted by that multiplicity. The Gaussian distribution describes how the represented force interacts with the carrier; it does not resolve the surface of a sphere.

The dimensionless fluid domain is $\Omega = [0, 3] \times [0, 1]$, periodic in x with stationary no-slip walls at $y = 0$ and $y = 1$. The channel has no open inlet or outlet. The initial fluid velocity is

$$\boldsymbol{u}_0(y) = (6y(1 - y), 0), \tag{5.1}$$

whose mean streamwise speed is one. A body-force acceleration $\boldsymbol{f}_d = (12/\mathrm{Re}_H, 0)$ maintains this profile in the unladen steady limit. Parcel initial velocities are zero. The reference depth ratio is $b = 1$, and the evaluated model uses $\mathrm{Re}_H = 20$, $R = 1000$, six parcels, and Gaussian width $\sigma = 0.14$.

5.2 Coupled equations

The fluid equations use the inertial pressure convention defined in equation (3.8):

$$\nabla \cdot \boldsymbol{u} = 0, \tag{5.2}$$

$$\partial_t \boldsymbol{u} + (\boldsymbol{u} \cdot \nabla)\boldsymbol{u} = -\nabla p_i + \frac{1}{\mathrm{Re}_H}\nabla^2 \boldsymbol{u} + \boldsymbol{S} + \boldsymbol{f}_d. \tag{5.3}$$

The parcel equations are

$$\dot{\boldsymbol{q}}_i = \boldsymbol{v}_i, \qquad \dot{\boldsymbol{v}}_i = \frac{\boldsymbol{U}_i - \boldsymbol{v}_i}{\mathrm{St}_i} + \boldsymbol{g}_{\mathrm{eff}}, \qquad \boldsymbol{F}_i = \frac{m_i}{\mathrm{St}_i}(\boldsymbol{U}_i - \boldsymbol{v}_i). \tag{5.4}$$

Here $\boldsymbol{g}_{\mathrm{eff}}$ is a specified effective acceleration, including any buoyancy adjustment in the physical mapping. The reaction source contains the hydrodynamic drag only.

Depositing parcel gravity again as fluid drag would count a different force twice.

5.3 Normalized Gaussian transfer

Let $\boldsymbol{x}_a$ denote the fixed transfer-quadrature locations and ΔA_a their area weights. A periodized raw Gaussian is evaluated between $\boldsymbol{x}_a$ and each parcel center. In the two-dimensional channel, its structure is

$$G_{ai} = \sum_{k\in\mathbb{Z}} \exp\left[-\frac{(x_a - q_{ix} + kL_x)^2 + (y_a - q_{iy})^2}{2\sigma^2}\right]. \tag{5.5}$$

The implementation retains a finite image set with distant tails beyond eight widths after choosing the nearest image. A global Gaussian prefactor is unnecessary because normalization removes it. Define

$$Z_i = \sum_a G_{ai}\Delta A_a, \qquad W_{ai} = G_{ai}/Z_i, \qquad \sum_a W_{ai}\Delta A_a = 1. \tag{5.6}$$

Both fluid interpolation and force spreading use these weights:

$$\boldsymbol{U}_i = \sum_a W_{ai}\boldsymbol{u}_a\Delta A_a, \qquad \boldsymbol{S}_a = -\sum_i W_{ai}\boldsymbol{F}_i. \tag{5.7}$$

Near a wall, Z_i accounts for the actual bounded-domain quadrature. At an arbitrary collocation point, the numerator uses that point, while the denominator remains the value from the fixed transfer quadrature.

5.3.1 Force and work identities

Substituting equation (5.7) into the quadrature sum gives

$$\sum_a \boldsymbol{S}_a\Delta A_a = -\sum_i \boldsymbol{F}_i, \tag{5.8}$$

and

$$\sum_a \boldsymbol{u}_a \cdot \boldsymbol{S}_a\Delta A_a = -\sum_i \boldsymbol{U}_i \cdot \boldsymbol{F}_i. \tag{5.9}$$

The combined interphase contribution to fluid and parcel kinetic-energy rates is therefore

$$\mathcal{P}_{\text{drag}} = \sum_i (\boldsymbol{v}_i - \boldsymbol{U}_i)\cdot\boldsymbol{F}_i = -\sum_i \frac{m_i}{\mathrm{St}_i}|\boldsymbol{U}_i - \boldsymbol{v}_i|^2 \le 0. \tag{5.10}$$

These are semidiscrete properties of the shared transfer. A trained PINN with nonzero residual need not satisfy the complete global momentum or energy balance exactly. External driving can also make total kinetic energy increase even though the drag

contribution is dissipative.

5.4 Neural representation and hard constraints

Two MLPs, each with three hidden layers of 32 tanh units, represent the fluid correction and parcel histories. The fluid network uses wall-normal position, time, and three periodic Fourier harmonics in x. The ansatz is

$$\boldsymbol{u}(\boldsymbol{x}, t) = \boldsymbol{u}_0(y) + \frac{t}{T} y(1-y)\boldsymbol{n}_u(\boldsymbol{x}, t). \tag{5.11}$$

It enforces the initial velocity and stationary wall velocity. Periodicity follows from the Fourier features. A time-dependent pressure offset is removed at $(0, 0)$; no independent initial-pressure condition is imposed.

The parcel network uses the relaxation time in its output envelopes:

$$b_i(t) = 1 - \exp(-t/\mathrm{St}_i), \qquad a_i(t) = t - \mathrm{St}_i b_i(t), \tag{5.12}$$

$$\boldsymbol{q}_i(t) = \boldsymbol{q}_{i0} + t\boldsymbol{v}_{i0} + a_i(t)\boldsymbol{n}_{q,i}(t), \qquad \boldsymbol{v}_i(t) = \boldsymbol{v}_{i0} + b_i(t)\boldsymbol{n}_{v,i}(t). \tag{5.13}$$

The envelope imposes the initial states and represents the known response scale. It supplies no target fluid velocity or target coupled trajectory. Position and velocity remain separate trainable functions, linked by $\dot{\boldsymbol{q}}_i - \boldsymbol{v}_i$ in the objective.

The weighted mean-square residual objective assigns weights 1 to fluid momentum, 2 to continuity, 5 to parcel kinematics, 1 to parcel acceleration, and 100 to center confinement. All calculations in this formulation use float64. Figure 5.1 emphasizes the common quadrature that couples the representations. The default calculations are collision-free; independently tested impact primitives are not treated as a validated collisional suspension model.

5.5 Reduced Stokes-response verification

The first reduced test, V01, solves the constant-forcing relaxation problem with dimensionless response time one:

$$\dot{v} = 1 - v, \qquad \dot{z} = v, \qquad v(0) = z(0) = 0. \tag{5.14}$$

Its analytical response is

$$v = 1 - e^{-t}, \qquad z = t - 1 + e^{-t}. \tag{5.15}$$

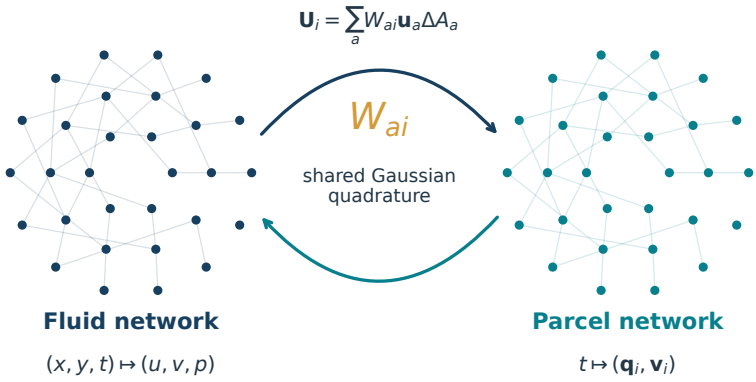


Figure 5.1. Coupled Gaussian formulation. The two stylized network symbols indicate distinct MLPs, each with three 32-unit hidden layers; their radial arrangement is illustrative and does not specify literal layer connectivity. Shared quadrature transfers sampled fluid velocity to parcels and hydrodynamic reaction to the fluid. The complete objective includes fluid and parcel residuals.

The reference arrays are excluded from training and used only after the physics-constrained fit. The relative velocity and displacement errors are 0.066% and 0.024%, with independent ODE residual RMS approximately 1.35×10^{-3}. Both the stated relative-error and residual criteria are satisfied. Figure 5.2 shows the stored predictions against the analytical response.

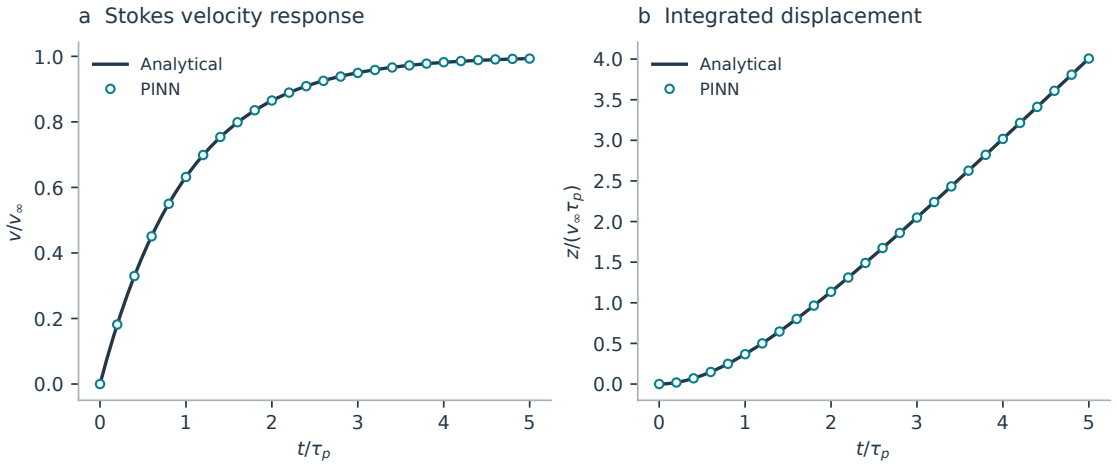


Figure 5.2. V01 reduced Stokes response. The PINN is fitted to the ODE residual and initial conditions; analytical values are held out. Open markers are subsampled stored predictions for legibility; the metrics use every evaluation sample. This test verifies the reduced response equation, not the full two-way channel.

5.6 Two-way momentum-exchange verification

The second reduced problem, V02, considers spatially uniform fluid and particle velocities, fluid mass one, parcel mass loading $\beta = 0.25$, and response time $\mathrm{St} = 0.2$. With $u(0) = 0$ and $v(0) = 1$,

$$\dot{u} = -\frac{\beta}{\mathrm{St}}(u - v), \qquad \dot{v} = \frac{u - v}{\mathrm{St}}. \tag{5.16}$$

The solution has constant total momentum $u + \beta v = 0.25$ and decaying slip. The relative fluid- and particle-velocity errors are 0.096% and 0.238%, respectively. The independent residual RMS is approximately 6.12×10^{-3}.

In this reduced PINN, the output relation $v = 1 - 4u$ imposes total momentum exactly. Its zero momentum error is consequently a hard-constraint check, not an independent demonstration that a general coupled solution conserves momentum. The nontrivial checks are the predicted relaxation and its ODE residual. Figure 5.3 also shows the kinetic-energy decay, consistent with the sign in equation (5.10).

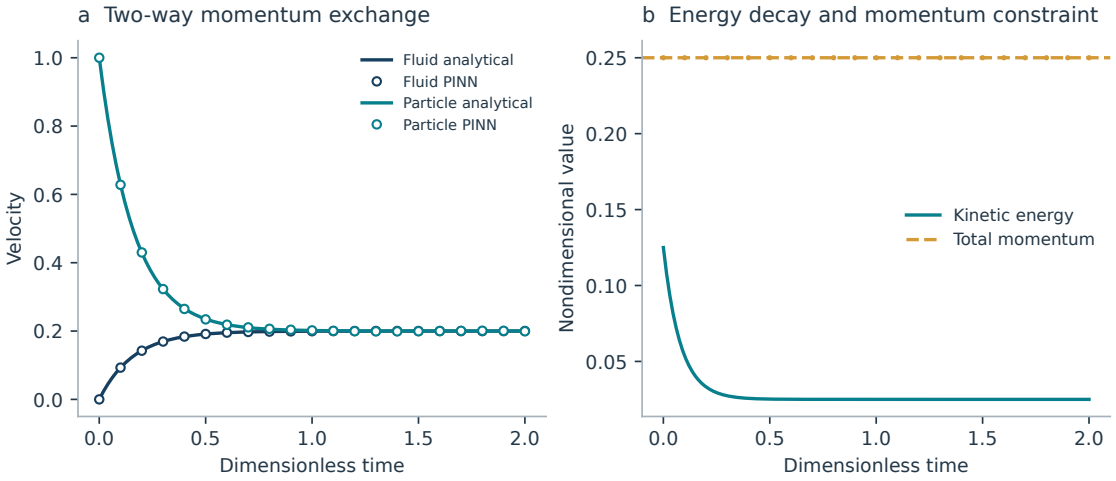


Figure 5.3. V02 uniform two-way relaxation. The two velocities approach a common value while kinetic energy decreases. The total-momentum line is exact by the imposed output relation. Velocity agreement and the independently sampled ODE residual assess whether the learned temporal response satisfies the governing problem.

5.7 Kernel and carrier-flow verification

5.7.1 Normalization, overlap, and parcel splitting

The bounded-domain normalization test V03 probes periodic-seam and wall locations. Its maximum recorded discrepancy is approximately 8.44×10^{-15} across the checks included in that test. The displayed integrals in figure 5.4 differ from one by only a few units of machine precision. The larger summary discrepancy also includes the constant-field sampling check and should not be assigned to an individual integral in the figure.

V04 compares a three-dimensional product quadrature with an independent free-space Gaussian overlap array. The relative difference is approximately 1.29×10^{-15}. The CSV column retained from the numerical workflow uses a generic prediction name; here it is correctly identified as a quadrature result rather than a learned PINN field. This is a kernel-integral test and does not establish a three-dimensional channel solution.

V09 splits a fixed population of 10,000 particles into 10, 50, and 100 coincident parcels. The maximum force-density difference is approximately 2.66×10^{-15}. Because the split centers coincide and total mass is unchanged, this tests deterministic representation invariance. It is not a convergence study of independently dispersed parcel populations.

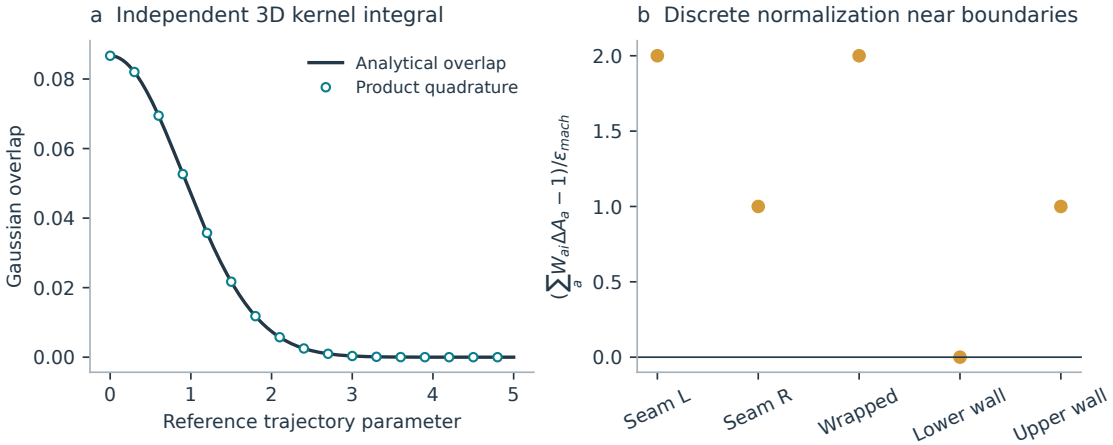


Figure 5.4. Gaussian transfer checks. Panel (a) shows V04's independent three-dimensional overlap integral and product quadrature. Panel (b) shows the stored V03 normalization values divided by machine precision. The kernels are verified on their declared quadrature; neither panel is a flow-field prediction.

5.7.2 Unsteady incompressible flow

V10 considers the periodic two-dimensional Taylor–Green vortex with viscosity $\nu = 0.01$ over $0 \le t \le 5$. The exact velocity is

$$u = \sin x \cos y \, e^{-2\nu t}, \qquad v = -\cos x \sin y \, e^{-2\nu t}, \tag{5.17}$$

and the gauge-aligned pressure is

$$p = \frac{1}{4}(\cos 2x + \cos 2y - 2)e^{-4\nu t}. \tag{5.18}$$

The initial field is imposed, while the analytical temporal decay is held out of training. Independent sample-relative errors are 0.218% for u, 0.187% for v, and 0.215% for pressure. The independent residual RMS is approximately 7.38×10^{-4}. These results support the unsteady carrier implementation for this smooth periodic test.

5.8 Status of the coupled straight-channel calculations

The evaluated coupled channel uses one training seed and four combinations of Stokes number and mass loading. The training schedule is 1,200 Adam steps followed by a 200-iteration L-BFGS budget; the transfer quadrature is 16×16. All four calculations produce finite outputs, but none meets the complete residual screen. Table 5.2 reports the momentum and parcel-acceleration RMS values that bound their interpretation.

The lowest listed momentum RMS is 0.0407, still above the declared bound. The full screen also requires all centers to remain inside the channel, a trained finite

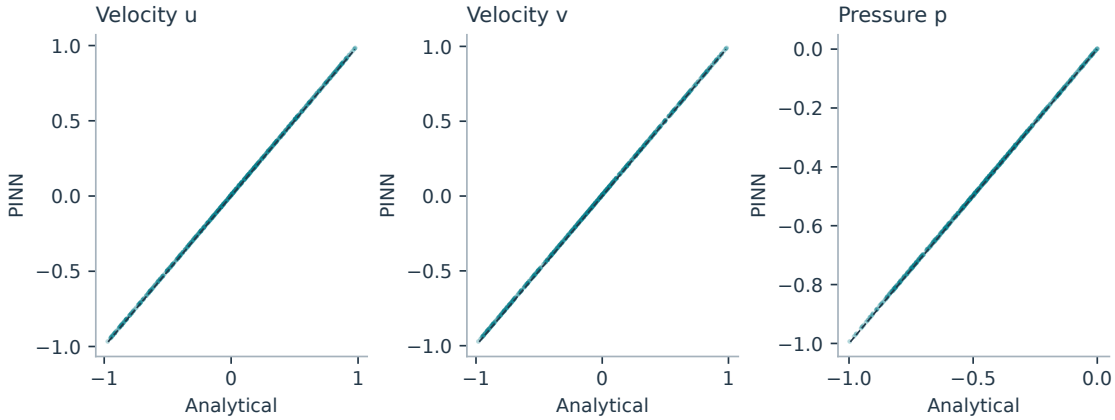


Figure 5.5. V10 Taylor–Green verification on independent space–time samples. Each panel compares the predicted quantity with its analytical value, using the same pressure gauge. The test contains no parcels and establishes carrier-flow verification rather than coupled particle-transport validation.

Table 5.1. Accepted solution-observable errors in the reduced Gaussian-model verification problems. Residual criteria are checked separately.

Test	Observable	Relative error
V01	Stokes velocity	0.066%
V01	Integrated displacement	0.024%
V02	Fluid velocity	0.096%
V02	Particle velocity	0.238%
V10	Taylor–Green u	0.218%
V10	Taylor–Green v	0.187%
V10	Taylor–Green p	0.215%

Table 5.2. Coupled straight-channel screening. Every required held-out residual must be below 0.03; the momentum residual exceeds that bound in all four cases. These are diagnostic calculations, not accepted transport predictions.

St	β	Momentum RMS	Parcel RMS	Complete screen
0.08	0.05	0.3854	0.0476	Not met
0.20	0.05	0.1372	0.0267	Not met
0.20	0.15	0.3526	0.0834	Not met
0.60	0.05	0.0407	0.0017	Not met

checkpoint without training error, local volume fraction below 0.01, and slip Reynolds number no greater than one. The observed maximum slip Reynolds numbers range from approximately 0.234 to 0.642. These satisfy that loose screening bound but do not establish the asymptotically small-slip regime required for accurate Stokes drag. Passing the reduced tests therefore does not justify selecting a visually smooth coupled-channel field as a validated result. I exclude these predictions from quantitative claims about loading effects, sedimentation, or deposition. The physical model also samples the disturbed fluid without an undisturbed-velocity correction, so successful future residual screening would still require independent assessment of the drag closure.

All ten component tests in the evaluated record pass their individual criteria. Tests V05–V08 concern impact algebra and a frozen-rate sampling primitive, which are outside the collision-free channel study and are not used to validate its dynamics. The literature catalogue likewise supplies no matched physical-validation result for the Gaussian channel. These scope limits preserve the distinction between tested components and an accepted coupled prediction.

5.9 Chapter conclusions

The Gaussian formulation establishes exact transfer-quadrature identities and successful reduced tests for relaxation, kernel normalization, and unsteady carrier flow. The strongest supported results are these independently evaluated components. Further optimization, transfer-grid and time-sampling refinement, multiple seeds, and a matched independent coupled reference are required before interpreting the fully coupled channel quantitatively. The next chapter therefore formulates a controlled constricted-channel extension rather than extending unsupported transport claims.

Chapter 6

Extension to a Constricted Channel

6.1 Purpose and status of the extension

A constriction introduces a varying gap, a changing cross-sectional mean velocity at fixed flux, and stronger sensitivity to wall treatment. It provides a natural next problem after the constant-width baseline and the Gaussian coupling study. It also introduces new requirements that cannot be met by merely changing the appearance of a straight-channel figure.

This chapter defines the geometry, governing problem, and evaluation design for that extension. The geometry drawings are mathematical illustrations. No constricted-channel velocity, pressure, particle trajectory, or deposition rate has been computed for the accepted result set. The proposed parameter values below are study-design choices rather than measured or completed results.

The obstruction is static and represented by the walls. A progressively growing particle deposit, moving obstacle, or porous plug would require additional evolution laws and is outside this design. I use the term *constriction* for the prescribed narrowing to keep this geometric effect separate from dynamic clogging.

6.2 Parameterized geometry

Let the nominal channel height be H and let the walls be $y = \pm h(x)$. A compact, smooth narrowing centered at x_c is defined by

$$h(x) = \frac{H}{2}[1 - \delta f(\xi)], \qquad \xi = \frac{x - x_c}{\ell}, \qquad 0 \le \delta < 1, \tag{6.1}$$

where

$$f(\xi) = \begin{cases} \left[\dfrac{1 + \cos(\pi\xi)}{2}\right]^2, & |\xi| < 1, \\ 0, & |\xi| \ge 1. \end{cases} \tag{6.2}$$

The profile and its first three derivatives match the constant-height segments at the joins. Its minimum gap is $H_t = (1-\delta)H$. The initial design uses $L_x/H = 8$, $x_c/H = 4$, and $\ell/H = 1.2$, with nominal severities $\delta \in \{0, 0.2, 0.4\}$. These values define a controlled sequence from the unblocked channel to moderate narrowing. The actual

reference study must check that upstream and downstream lengths are sufficient.

Figure 6.1 displays the geometry and height profiles. Only geometric quantities are plotted. In particular, the height curves must not be interpreted as velocity or pressure results.

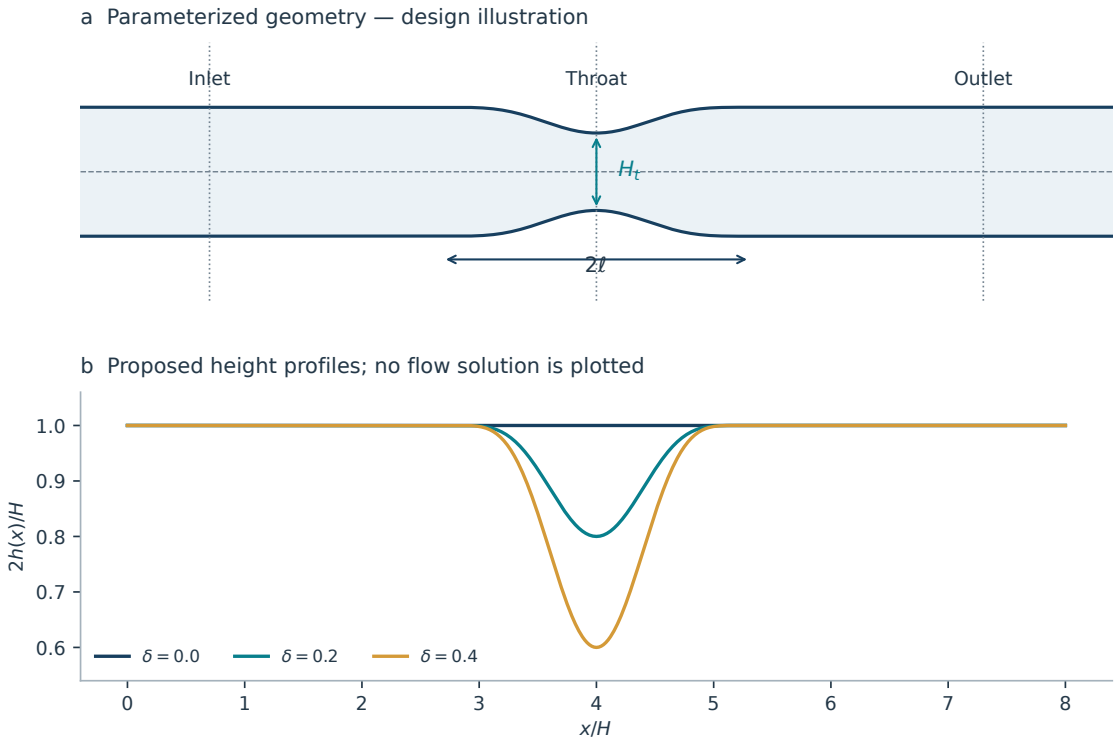


Figure 6.1. Proposed constriction. Panel (a) defines the throat, characteristic length, and reference sections. Panel (b) shows the mathematical height profiles for the proposed severity values. The illustrated channel has $L_x/H = 8$ and $\ell/H = 1.2$; no fluid or particle solution is represented.

6.3 Carrier-flow boundary conditions

For an open-channel extension, the governing equations retain the inertial pressure scale and the same Newtonian constitutive law. I prescribe a parabolic inlet velocity with a specified flux Q and no slip on both walls. A downstream traction condition provides an outlet reference rather than imposing a second velocity profile. The finite-element reference and the PINN must implement exactly the same outlet condition and pressure convention.

This is a different boundary-value problem from both earlier stages. The curved baseline has prescribed velocities at both ends, whereas the straight Gaussian channel is periodic and body-force driven. Comparisons across these stages must therefore identify which changes follow from geometry and which follow from the drive or

boundary conditions. A useful control is a straight open channel with the same inlet and outlet treatment as the constricted case.

For unsteady particle introduction, the initial carrier field should be an independently converged unladen solution in the constricted domain. Imposing the original straight-channel profile throughout the narrowed region would violate the new wall geometry. An alternative start from rest requires a compatible smooth inlet ramp and its own transient verification. The primary design uses the steady unladen field to avoid mixing inlet startup with particle response.

6.4 Geometry-aware representation

The fixed-width centerline mapping of the baseline cannot represent this variable-height domain without modification. A convenient mapping is

$$(x, \eta) \longmapsto (x, y) = (x, h(x)\eta), \qquad -1 \leq \eta \leq 1. \tag{6.3}$$

Its area Jacobian is $h(x)$. If the differential residual is evaluated in mapped coordinates, the chain rule must include this geometry dependence. For example,

$$\partial_x|_y = \partial_x|_\eta - \eta \frac{h'(x)}{h(x)} \partial_\eta, \qquad \partial_y = \frac{1}{h(x)} \partial_\eta. \tag{6.4}$$

Evaluating automatic derivatives directly with respect to physical coordinates is another option. Either route requires a manufactured derivative test before training a new flow problem.

Periodic Fourier features in the streamwise direction would impose the wrong structure for the proposed open domain. They must be replaced by nonperiodic coordinate features. A wall-distance factor can enforce the stationary wall velocity, but it must be consistent with the nonzero inlet condition and with the mapped derivatives. A divergence-free streamfunction ansatz is a possible alternative, provided that its boundary and pressure conditions are verified. These are implementation choices to evaluate, rather than capabilities already established by the straight-channel model.

6.5 Gaussian support near the narrowing

The force-transfer rule must use quadrature over the fluid domain:

$$Z_i = \sum_{a \in \Omega_\delta} G(\boldsymbol{x}_a - \boldsymbol{q}_i) \Delta A_a, \qquad W_{ai} = \frac{G(\boldsymbol{x}_a - \boldsymbol{q}_i)}{Z_i}. \tag{6.5}$$

For mapped quadrature, ΔA_a includes the Jacobian from equation (6.3). The same points, weights, and normalizer must be used for interpolation and reaction-force deposition. Under that condition, the derivations of equations (5.8) and (5.9) remain valid on the new quadrature.

Figure 6.2 illustrates why normalization must follow the actual domain. Some Gaussian support intersects the wall region near the throat. Removing those points without updating Z_i would lose force. Renormalizing restores the integrated force, but changes the shape of the sampling operator. It still does not supply missing lubrication or undisturbed-velocity corrections.

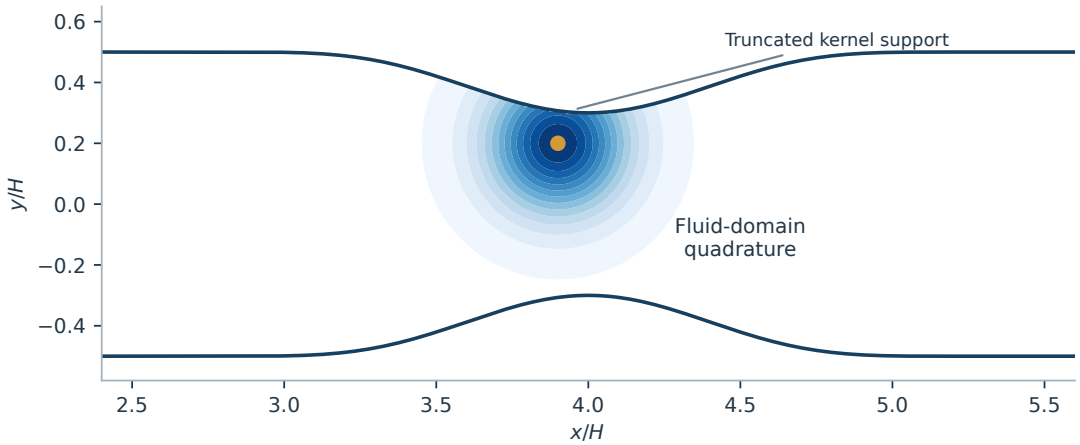


Figure 6.2. Illustration of a Gaussian centered near the upper throat wall. Shading is a prescribed kernel shape restricted to the proposed fluid domain, not a computed concentration or force field. The normalizer must use the same fluid-domain quadrature as both transfer operations.

Three ratios require explicit monitoring:

$$\frac{d}{H_t}, \qquad \frac{\sigma}{H_t}, \qquad \frac{\sigma}{\Delta_{\text{quad}}}, \tag{6.6}$$

where Δ_{quad} denotes a local quadrature-spacing measure. The first concerns physical confinement; the second concerns regularization relative to the gap; the third concerns the numerical resolution of that regularization. They address different questions and should be varied separately. If physical wall effects dominate, the present unbounded Stokes closure may need replacement before more accurate training is useful.

6.6 Particle and parcel wall conditions

An individual sphere has an admissible center region obtained by offsetting the wall by its radius. For a sloped wall, the offset is based on signed normal distance to the boundary, not only the vertical gap. A contact event must therefore use the actual

geometry. The baseline's irreversible sticking rule can be retained as a controlled modeling choice, but it must not be called a prediction of deposition physics.

A parcel introduces an additional issue: its center represents many particles. Stopping the entire represented population at the first center contact may create a strong and artificial capture rule. The proposed first study therefore focuses on carrier fields and freely transported dilute parcels whose centers remain inside the admissible region. A deposition extension should specify whether it resolves individual particles, splits parcels, or uses a separately justified population-level wall model. No capture efficiency is assigned before making that choice.

6.7 Controlled evaluation sequence

The extension should be evaluated in the following order.

1. Recover the straight, unladen open-channel solution at $\delta = 0$ with the new inlet and outlet conditions. Verify pressure gauge, wall velocity, and section flux before adding particles.
2. Solve the unladen constricted channel with an independent mesh-based method. Refine the throat and the upstream and downstream transition regions; monitor pressure drop and the complete velocity field.
3. Introduce individual one-way particles with identical initial conditions in the neural and reference fields. Compare position, common-time coverage, and endpoint class separately.
4. Add a single Gaussian parcel at low loading and verify force/work transfer on the irregular-domain quadrature. Compare the coupled fluid response with an independently discretized regularized model.
5. Vary severity and loading only after the earlier controls pass. Keep physical diameter, total represented mass, and kernel width explicit.

The initial proposed Reynolds numbers are 1, 5, and 20 at nominal height. These are design points rather than an observed regime map. Particle response times should be selected to retain an acceptably small slip Reynolds number at the throat. Mass-loading changes must be made through represented population while holding diameter fixed when the purpose is to isolate feedback. A separate diameter sweep changes both response time and confinement, and should be interpreted accordingly.

6.8 Observables and acceptance criteria

The primary observables are the section flux, throat velocity profile, pressure drop between fixed reference sections, and field error against the matched reference. For particles, residence time and endpoint class are added only after trajectory agreement is demonstrated. The reference sections must remain identical across severity values so that the pressure drop represents the same geometric interval.

At fixed flux, incompressibility gives the exact sectional mean relation

$$\overline{u}_x(x) = \frac{Q}{2h(x)}, \qquad \frac{\overline{u}_{x,t}}{\overline{u}_{x,0}} = \frac{1}{1-\delta}. \tag{6.7}$$

This is a geometric conservation implication, not a computed centerline velocity or a prediction of recirculation. A lubrication estimate, derived in chapter A, gives a useful additional check only when its small-slope and low-inertia assumptions apply.

Table 6.1. Proposed acceptance protocol for the constricted study. Limits must be fixed before interpreting the associated parameter sweep; these checks have not been executed.

Quantity	Proposed check	Interpretation
Unladen field	$E_u, E_p < 5\%$ against a matched reference	Carrier-field agreement
Reference refinement	Changes in Δp and field norms below 2%	Numerical sensitivity, not exact uncertainty
Section flux	Maximum relative variation below 2%	Integrated mass balance
Transfer	Force/work discrepancies consistent with roundoff	Algebraic operator consistency
Particle paths	RMS position below $0.05W$; common interval above 95%; all endpoint classes matched	Transport accuracy for stated initial conditions
Physical regime	Report maximum Re_p, d/H_t, local α and wall distance	Closure applicability; no universal threshold inferred

The error limits in table 6.1 are proposed numerical screens. In addition, residuals should be reported by component on a separate set of throat, near-wall, and bulk points. A global average must not conceal a concentration of error near the narrowing. The Gaussian width study should remain separate from spatial and temporal refinement, and multiple seeds should be included before drawing a quantitative trend.

6.9 Expected mechanisms and testable statements

Conservation at fixed flux predicts a larger section-mean speed as the gap narrows. The accompanying changes in particle transit depend on response time, local flow, and the retained force law. I treat changes in residence time, wall approach, and downstream recovery as hypotheses for the controlled study. Recirculation, enhanced deposition, or a particular pressure-loss factor cannot be inferred from the geometry alone.

At fixed imposed pressure difference, flux becomes an outcome instead. That experiment addresses a different question and must not be pooled with the fixed-flux design. Keeping these operating conditions separate is essential for a meaningful comparison of unblocked and constricted cases.

6.10 Chapter conclusions

The constricted-channel stage is fully specified at the level of geometry, governing conditions, transfer consistency, and evaluation requirements. It extends the study in a controlled direction while preserving the distinction between mathematical consequences, proposed experiments, and completed numerical findings. Its next milestone is a verified unladen reference problem, followed by one-way transport and then two-way Gaussian feedback.

Chapter 7

Discussion

7.1 What the baseline establishes

The baseline demonstrates that a physics-informed field representation can agree closely with an independently discretized curved-channel solution for the specified steady, low-Reynolds-number problem. The velocity and pressure errors are below 2% and 1%, respectively, while the drag-network error is below 1%. The analytical plane-flow test provides an additional check that is independent of the curved-channel finite-element field. Reference-mesh sensitivity and section flux add evidence beyond a single parity plot.

The trajectory comparison changes the interpretation of this success. The carrier-field results are useful but do not establish that the complete particle problem is solved with the same relative accuracy. A particle integrates the local field over time, and a small near-wall discrepancy can change an endpoint event. Six of seven matching endpoint classes therefore indicate a remaining limitation rather than permission to generalize deposition trends. This finding supports evaluating the final transport quantity directly, instead of treating bulk velocity error as its surrogate.

The force network also raises a methodological question. Since the retained Stokes law is known analytically, its neural approximation is most useful as a controlled component of a broader learning framework. The present error does not show that the learned force is superior to the direct analytical formula. A future assessment should compare the two within identical carrier fields and identify whether the approximation adds cost, flexibility, or error.

7.2 What conservative transfer establishes

The shared Gaussian operator gives a stronger result of a different kind. The force and work identities follow from the normalization and the common quadrature. Their validity does not depend on a low training loss. This property removes one avoidable source of inconsistent interphase exchange and remains available when the carrier geometry changes, provided that the same fluid-domain quadrature is used consistently.

The identities do not guarantee the accuracy of the coupled solution. The carrier momentum equation can retain a residual even when the sum of deposited forces is

exact. In the evaluated coupled channel, the remaining momentum residual exceeds the study's bound in every tested parameter combination. The correct inference is that the operator design is verified while the quantitative coupled-channel prediction remains unresolved.

The successful reduced tests address important parts of the calculation: response-time dynamics, reciprocal momentum exchange, kernel integrals, and unsteady incompressible flow. Each isolates a necessary capability. Their combination is a verification chain, not a substitute for a matched independent two-way channel reference. This distinction is consistent with the separation between numerical verification and physical validation [11].

7.3 Optimization, residuals, and independent error

The completed calculations illustrate why optimization metrics and independent solution errors should be reported together. Residual normalization changes their numerical magnitude, and a weighted objective can favor some components over others. Reported raw residuals should therefore retain their equation convention. In particular, comparing the baseline momentum RMS directly with the Gaussian momentum RMS would ignore different scaling, geometry, and transient content.

Training difficulties in physics-informed learning are well documented [3, 4]. They motivate focused improvements such as component scaling, more informative collocation, and better temporal resolution of the response scale. These improvements must be assessed using the same independent checks. Lowering a threshold after inspecting the results or selecting only a visually attractive field would weaken the inference rather than improve the model.

7.4 Regularization and model validity

The Gaussian width controls how the represented force is distributed. It affects both the sampled fluid velocity and the carrier response. A smaller width is not automatically more physically accurate, especially if the transfer quadrature underresolves it. A larger width can smooth localized forcing while changing the regularized interaction. Width sensitivity must therefore be distinguished from convergence at fixed model parameters.

The undisturbed-velocity issue is equally important. The current parcels sample the disturbed fluid directly. A reduced Stokes-response test cannot resolve the resulting self-interaction question in a spatially coupled flow. Comparison with an independent solver of the same regularized equations would test numerical consistency; assessment

against a suitable physical reference would additionally test the closure.

Confinement further separates numerical and physical accuracy. A model can converge accurately to equations that omit important wall effects. For that reason, the proposed constriction study records particle-to-gap ratio and local slip Reynolds number in addition to ordinary solver errors. If those quantities indicate a breakdown of the retained approximation, more training is not the appropriate physical correction.

7.5 Transferability across the three stages

Table 7.1 summarizes the actual scope of the findings. The baseline establishes a carrier-field comparison for one geometry and Reynolds number, with two seeds. The Gaussian tests establish specific reduced dynamics and transfer properties. The constriction stage provides a defined study design. None of these establishes a universal neural operator over channel shape, loading, and Reynolds number.

Table 7.1. Supported conclusions and their boundaries.

Stage	Supported statement	Remaining boundary
Baseline carrier	Analytical and independent numerical field agreement	One declared geometry/Re case; finite residuals
Baseline particles	Completed integrations and drag approximation	Endpoint and full-path criteria not all met
Gaussian transfer	Discrete force/work consistency and component verification	Disturbed-velocity closure and coupled field error
Straight coupled channel	Executed diagnostic predictions	Full residual screen not met
Constriction	Defined geometry and evaluation protocol	No completed constricted-channel solution

The transferable parts are the nondimensional bookkeeping, reference matching, shared transfer, and evaluation logic. New physical and geometric claims require their own executed evidence. In particular, a parameter slot in a network does not establish learned parametric generalization, and a geometric illustration does not establish a flow solution.

7.6 Implications for the next computational study

The most informative next work is narrower than a large parameter sweep. First, the baseline trajectory limitation should be isolated using the same carrier field with exact and learned drag. Second, the Gaussian channel should be compared with an

independent regularized solver at fixed width and well-resolved quadrature. Third, the unladen constricted problem should be solved before introducing moving populations.

This ordering reduces ambiguity. A failed two-way constricted calculation could otherwise reflect carrier approximation, irregular-domain transfer, wall geometry, particle closure, or optimization. Establishing the controls in sequence allows each additional mechanism to be interpreted against a known numerical foundation.

Chapter 8

Conclusions and Future Work

8.1 Principal conclusions

This thesis develops a staged framework for physics-informed fluid–particle transport. The progression begins with a one-way coupled particle problem, introduces Gaussian parcel reaction in a straight channel, and defines a smooth constriction for the next study. The completed evidence supports the following conclusions.

First, the baseline carrier network agrees with an independent finite-element solution of the stated curved-channel problem, with velocity and gauge-aligned pressure errors of 1.510% and 0.863%. Analytical plane-flow verification and reference-mesh sensitivity support this comparison. The normalized learned Stokes response has an independent error of 0.717%.

Second, carrier-field agreement does not by itself establish accurate particle transport. The seven baseline comparisons do not all meet the endpoint, position, and time-coverage conditions. Deposition and wall-sensitive parameter trends are therefore outside the accepted quantitative conclusions.

Third, a normalized Gaussian used for both interpolation and spreading preserves discrete force and work on the transfer quadrature. The resulting interphase drag power is nonpositive. These identities are algebraic properties of the discretization and do not imply that an imperfectly trained field satisfies every global balance.

Fourth, the reduced Gaussian-model tests recover Stokes response, uniform two-way relaxation, kernel integrals, and Taylor–Green decay with the reported errors and residuals. The accepted solution-observable errors in the listed relaxation and vortex tests are below 0.25%. The fully coupled straight-channel predictions remain diagnostic because their complete residual criteria are not met.

Finally, the constricted-channel extension has a consistent geometric and evaluation definition, including physical-coordinate derivatives, fluid-domain kernel normalization, and independent throat-resolution checks. It has no completed predictive result in the present study.

8.2 Contributions

The study contributes a common evaluation framework connecting carrier fields, drag approximation, transfer consistency, and particle outcomes. It gives a dimensionally consistent planar sphere-parcel formulation and explicitly separates a Gaussian force distribution from a resolved solid body. It also defines which verified components can be retained when extending the model to a variable-height channel.

These contributions provide a practical basis for further computation without assigning unsupported accuracy to complex coupled predictions. The numerical record and portable manuscript sources preserve the arrays, criteria, and figure provenance needed to reassess the stated findings.

8.3 Priorities for future work

The immediate priority is to improve and independently verify the transport-sensitive parts of the calculation. For the baseline, this means near-wall field assessment and a direct exact-drag comparison. For Gaussian coupling, it means a matched independent regularized channel solution, transfer-quadrature refinement, temporal sampling refinement, and multiple training seeds at fixed physical parameters.

The constricted study should then begin with an unladen flow and a straight-channel boundary-condition control. One-way particles can be introduced after the carrier reference passes, followed by low-loading Gaussian feedback. Physical deposition or dynamic clogging requires a separately specified wall and population model. Experimental comparison will require matched geometry, operating conditions, particle properties, and observables before it can establish physical validity.

References

[1] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019. doi:10.1016/j.jcp.2018.10.045.

[2] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, and L. Yang, "Physics-informed machine learning," *Nature Reviews Physics*, vol. 3, pp. 422–440, 2021. doi:10.1038/s42254-021-00314-5.

[3] S. Wang, Y. Teng, and P. Perdikaris, "Understanding and mitigating gradient flow pathologies in physics-informed neural networks," *SIAM Journal on Scientific Computing*, vol. 43, no. 5, pp. A3055–A3081, 2021. doi:10.1137/20M1318043.

[4] S. Wang, X. Yu, and P. Perdikaris, "When and why PINNs fail to train: A neural tangent kernel perspective," *Journal of Computational Physics*, vol. 449, article 110768, 2022. doi:10.1016/j.jcp.2021.110768.

[5] M. R. Maxey and J. J. Riley, "Equation of motion for a small rigid sphere in a nonuniform flow," *The Physics of Fluids*, vol. 26, pp. 883–889, 1983. doi:10.1063/1.864230.

[6] R. Clift, J. R. Grace, and M. E. Weber, *Bubbles, Drops, and Particles*. New York: Academic Press, 1978.

[7] P. Gualtieri, F. Picano, G. Sardina, and C. M. Casciola, "Exact regularized point particle method for multiphase flows in the two-way coupling regime," *Journal of Fluid Mechanics*, vol. 773, pp. 520–561, 2015. doi:10.1017/jfm.2015.258. Preprint: `https://arxiv.org/abs/1405.6969`.

[8] M. Uhlmann, "An immersed boundary method with direct forcing for the simulation of particulate flows," *Journal of Computational Physics*, vol. 209, no. 2, pp. 448–476, 2005. doi:10.1016/j.jcp.2005.03.017.

[9] C. Taylor and P. Hood, "A numerical solution of the Navier–Stokes equations using the finite element technique," *Computers & Fluids*, vol. 1, no. 1, pp. 73–100, 1973. doi:10.1016/0045-7930(73)90027-3.

[10] T. Gustafsson and G. D. McBain, "scikit-fem: A Python package for finite element assembly," *Journal of Open Source Software*, vol. 5, no. 52, article 2369, 2020. doi:10.21105/joss.02369.

[11] P. J. Roache, *Verification and Validation in Computational Science and Engineering*. Albuquerque, New Mexico: Hermosa Publishers, 1998.

[12] J. H. Ferziger and M. Perić, *Computational Methods for Fluid Dynamics*, 3rd ed. Berlin: Springer, 2002. doi:10.1007/978-3-642-56026-2.

[13] G. K. Batchelor, *An Introduction to Fluid Dynamics*. Cambridge: Cambridge University Press, 1967.

[14] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," in *Proceedings of the 3rd International Conference on Learning Representations*, 2015. `https://arxiv.org/abs/1412.6980`.

[15] D. C. Liu and J. Nocedal, "On the limited memory BFGS method for large scale optimization," *Mathematical Programming*, vol. 45, pp. 503–528, 1989. doi:10.1007/BF01589116.

Appendix A

Derivations and Transport Limits

A.1 Uniform two-way relaxation

For the uniform system in equation (5.16), define slip $w = v - u$. Adding the fluid and population-weighted particle equations gives

$$\frac{d}{dt}(u + \beta v) = 0. \tag{A.1}$$

Subtracting the equations gives

$$\dot{w} = -\frac{1+\beta}{\mathrm{St}} w, \qquad w(t) = w_0 \exp\left[-\frac{1+\beta}{\mathrm{St}} t\right]. \tag{A.2}$$

If $M_0 = u_0 + \beta v_0$ is the conserved momentum, the velocities are

$$u(t) = \frac{M_0 - \beta w(t)}{1+\beta}, \qquad v(t) = \frac{M_0 + w(t)}{1+\beta}. \tag{A.3}$$

For $u_0 = 0$, $v_0 = 1$ and $\beta = 0.25$, both velocities tend to 0.2. The total kinetic energy is

$$K = \frac{1}{2}u^2 + \frac{\beta}{2}v^2, \qquad \dot{K} = -\frac{\beta}{\mathrm{St}}(v-u)^2. \tag{A.4}$$

This gives the dissipative reference behavior of V02. It concerns a closed, uniform exchange system; the driven channel has additional work and boundary terms.

A.2 Adjoint transfer and kinetic energy

Let J denote the interpolation operator with entries $J_{ia} = W_{ai}\Delta A_a$. With quadrature mass matrix $M_A = \mathrm{diag}(\Delta A_a)$, the spreading operator is $-M_A^{-1}J^{\mathsf{T}}$ when the parcel force vector stores integrated forces. Thus

$$\boldsymbol{U} = J\boldsymbol{u}, \qquad \boldsymbol{S} = -M_A^{-1}J^{\mathsf{T}}\boldsymbol{F}. \tag{A.5}$$

The fluid work is $\boldsymbol{u}^{\mathsf{T}} M_A \boldsymbol{S} = -(J\boldsymbol{u})^{\mathsf{T}}\boldsymbol{F}$. Adding parcel drag work produces equation (5.10). This matrix statement makes explicit why interpolation and spreading must use the same quadrature and weights.

For a collision-free driven channel with no-slip stationary walls and periodic streamwise faces, a formally exact solution obeys

$$\frac{d}{dt}\left[\frac{1}{2}\int_\Omega |\boldsymbol{u}|^2\,dA + \sum_i \frac{m_i}{2}|\boldsymbol{v}_i|^2\right] = \int_\Omega \boldsymbol{u}\cdot\boldsymbol{f}_d\,dA + \sum_i m_i \boldsymbol{v}_i\cdot\boldsymbol{g}_{\text{eff}} - \frac{1}{\mathrm{Re}_H}\int_\Omega |\nabla\boldsymbol{u}|^2\,dA - \sum_i \frac{m_i}{\mathrm{St}_i}|\boldsymbol{U}_i-\boldsymbol{v}_i|^2. \tag{A.6}$$

Pressure and periodic-face contributions cancel under the stated conditions. The wall velocity is zero, so stationary walls do no work. Quadrature and neural equation residuals can create a nonzero discrepancy when this identity is evaluated numerically. Wall tractions still enter the momentum budget; they must not be omitted merely because wall work vanishes in the energy budget.

A.3 A limiting pressure-drop estimate

For slowly varying plane walls and negligible inertia, the local lubrication approximation gives the dimensional velocity

$$u(x,y) = -\frac{1}{2\mu}\frac{dp}{dx}[h(x)^2 - y^2]. \tag{A.7}$$

Integrating over $-h \le y \le h$ yields

$$Q = -\frac{2h(x)^3}{3\mu}\frac{dp}{dx}, \qquad \Delta p = \frac{3\mu Q}{2}\int_{x_a}^{x_b}\frac{dx}{h(x)^3}. \tag{A.8}$$

For constant half-height h_0, this reduces to the plane Poiseuille pressure drop $3\mu Q(x_b - x_a)/(2h_0^3)$. At fixed flux, the normalized hydraulic resistance estimate is

$$\frac{\Delta p}{\Delta p_0} = \frac{h_0^3}{x_b - x_a}\int_{x_a}^{x_b}\frac{dx}{h(x)^3}. \tag{A.9}$$

This is a limiting analytical check, not a simulation result for the proposed severity values. Its reliability requires sufficiently small wall slope and inertial effects; a full Navier–Stokes solution is needed when those assumptions fail [13].

A.4 Baseline trajectory criteria

Table A.1 retains the seven baseline comparisons associated with the field results in chapter 4. RMS position discrepancy is normalized by $W = 0.3$. The current numerical limits require discrepancy below 5%, common-time coverage above 95%,

and agreement of every endpoint class. Some individual rows satisfy those conditions, but the complete set does not. The trajectories are consequently not used as validated deposition evidence.

Table A.1. Complete baseline particle comparison; no initial position is removed to improve the aggregate outcome.

Initial η	RMS position/W (%)	Common time (%)	Same endpoint class
-0.650	2.67	98.86	Yes
-0.433	2.30	99.33	Yes
-0.217	4.30	99.27	Yes
0.000	7.83	98.59	Yes
0.217	16.09	96.19	Yes
0.433	34.97	87.39	Yes
0.650	3.29	73.55	No

A.5 Planar mass convention

For a physical sphere of diameter d, mass is $\rho_p \pi d^3/6$. A parcel containing N_i such spheres has mass $N_i \rho_p \pi d^3/6$. Dividing by $\rho_f H^2 B$ gives

$$\frac{N_i \rho_p \pi d^3/6}{\rho_f H^2 B} = \frac{N_i R \pi (d/H)^3}{6(B/H)}, \tag{A.10}$$

which is equation (3.9). The dimensional reaction source in a planar fluid equation has force divided by depth and area. This is why a two-dimensional Gaussian normalization uses area rather than volume. The reference depth records the physical interpretation without introducing a resolved spanwise flow field.

Appendix B

Computational and Reproducibility Record

B.1 Completed numerical records

The baseline results are associated with the computational record `basic_particle_problem/results/final_study`. The Gaussian verification results and the diagnostic coupled calculations are associated with `final_cloud_code/outputs/evaluated_2d_final`. These two result sets are kept separate from short environment smoke tests. The smoke tests are not used as research results in this thesis.

The archived repository revisions are

```
basic_particle_problem:
c3eb0efe76bcd62a3299dcd6c28922dbfa4c667a

final_cloud_code:
9d996c4f555bf5bf09e9f6fa78a3877e44310ec9
```

The source bundle contains a SHA-256 manifest of the numerical inputs and method files used to prepare the manuscript. The original result files are copied without numerical modification. Figure scripts change presentation, not the stored solution values. Gauge alignment and the exact computations used to recompute metrics are explicit in those scripts.

The historical result manifests retain their own source hashes. The archived baseline model and geometry, and the Gaussian model, physics, and verification modules match their recorded hashes. Current driver and configuration files do not all match the historical manifests; they are retained for method inspection and are not claimed as an exact replay of the original training execution.

B.2 Training and evaluation distinction

The baseline final evaluation restores previously trained network weights. The associated original flow configuration uses 1,000 Adam epochs, a 250-iteration L-BFGS budget, and a refinement budget of 1,800 epochs with 700 L-BFGS iterations. The

original force budget is 1,000 epochs. The two flow seeds are 42 and 43. The later evaluation configuration contains different default training fields, which do not imply that restored weights were retrained with those defaults.

The Gaussian channel uses seed 7, float64, and the evaluated configuration with 1,200 fluid/parcel Adam steps and a 200-iteration L-BFGS budget. Its reduced ODE tests use 700 Adam steps with a 120-iteration refinement budget; the Taylor–Green test uses 1,200 Adam steps with a 200-iteration refinement budget. An iteration budget need not equal the number of closure evaluations in a line-search optimizer. The stored status files and configuration distinguish these quantities.

Reference values are excluded from the Gaussian reduced-test optimization. In V01 the initial response is imposed. In V02 the initial values and total momentum relation are imposed. In V10 the initial velocity and pressure structure are imposed while the temporal decay is held out. These choices determine which aspects of the tests are independent predictions.

B.3 Software and execution scope

The numerical environment is Linux with Python 3.12. The baseline CPU environment uses PyTorch 2.14.1, and the Gaussian CPU environment uses PyTorch 2.6.0. The scientific stack includes NumPy 2.3.5, SciPy 1.17.0, Matplotlib 3.10.8, and scikit-fem 12.0.2. The preserved historical Gaussian runtime file records CPU execution with two threads. These are computational reproducibility details; no GPU performance or speedup is claimed.

The source bundle supports two levels of reproduction. Compiling `thesis.tex` uses the included tables and vector figures and does not require retraining. Rebuilding the figures additionally uses the archived arrays, model definitions, and two small baseline checkpoints. Independent metric checks include the full Poiseuille velocity vector, centered pressure, the final held-out drag evaluation, and the reduced Gaussian benchmarks.

Reproducing all training runs requires the complete corresponding repositories and their dependency environments. It also requires their reference-input files. The manuscript bundle is a reproducible evidence and typesetting package, not a claim that every full simulation can be retrained from the manuscript alone.

B.4 Figures and numerical transformations

Fourteen figures are included, each registered with its input source or its schematic status. Architecture figures are explanatory drawings. The constriction and kernel-wall

illustrations specify design choices. They are never registered as computed flow results. Numerical plots retain all evaluation arrays for aggregate errors; marker subsampling is used only to make reference curves readable.

Pressure comparisons subtract an additive constant independently from each array. Paired field panels use common color limits. Discrepancy panels have their own explicitly described scale and are not smoothed to conceal local error. The drag parity plot uses a stored diagnostic set different from the final held-out metric set; both are identified in the caption and validation record.

No failed channel field, synthetic demonstration, or extruded three-dimensional illustration is used as accepted predictive evidence. The material limitations are retained in the main chapters and the complete baseline trajectory table. Additional unsuccessful run histories remain outside the figure selection because they do not add accepted transport results.

B.5 Data and availability record

The computational projects are `https://github.com/mohrezsayeh90-ai/basic_particle_problem` and `https://github.com/mohrezsayeh90-ai/final_cloud_code`. The accompanying manuscript source package contains the selected numerical evidence, figure-generation scripts, numerical audit, bibliographic entries, and complete LaTeX source. It contains no experimental measurements produced by this study and no constricted-channel solution data. External literature is cited for its methods and physical context; it is not represented as measured agreement with the present channel model.